

[544] Zygotes

Tyler Caraza-Harter

Outline


Zygote Background

Zygote Limitations


Sandboxed Zygotes

OpenLambda Zygotes: Policy

Single-Threaded Process




Single-Threaded Process



do some initialization (e.g., importing a module):
uses time and space

Single-Threaded Process




fork() creates a new process.

(note: fork is a wrapper around clone, which takes more args)

The two processes (parent and child) will have (initially) identical address spaces, backed by the same physical memory

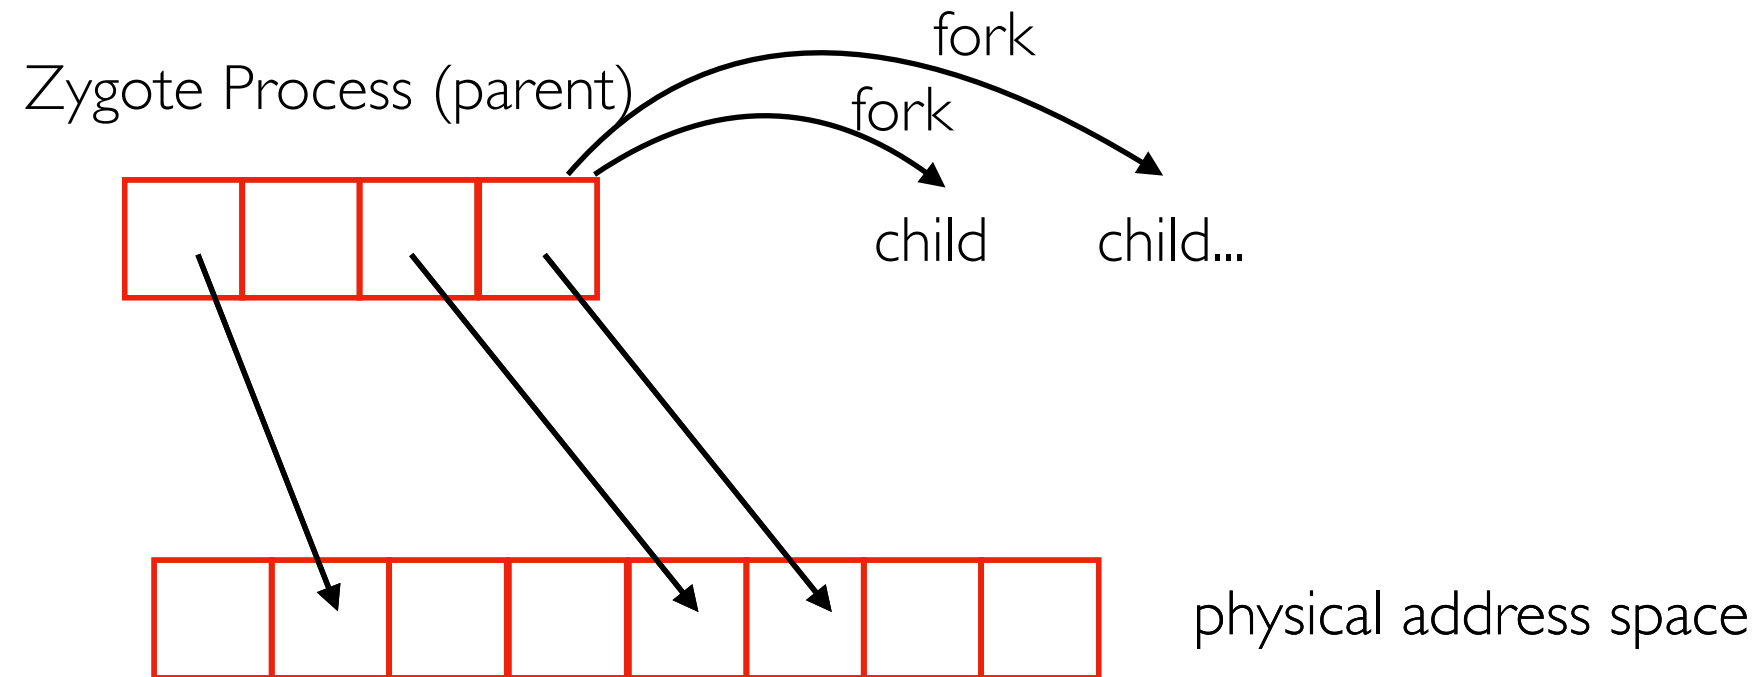
Single-Threaded Process



copy on write (COW):

if either of the processes write a page, a copy is made (so as not to interfere with the other process)


Zygotes




Zygote definition: a pre-initialized process that exists solely for the purpose of efficiently (with respect to space and time) forking new processes

Flat vs. Hierarchical

flat




hierarchical




Zygotes in Android (first use?)

Enter The Zygote



- nascent VM process
- starts at boot time
- preloads and preinitializes classes
- **fork()**s on command



slide from 2008 Google I/O talk: Dalvik VM Internals

<https://sites.google.com/site/io/dalvik-vm-internals>

Outline

Zygote Background

Zygote Limitations

Sandboxed Zygotes

OpenLambda Zygotes: Policy


Issue 1: Address Space Randomization

Exploiting security bugs is easier if you know where things are (e.g., injecting a function addr as part of a buffer overflow).

Address space randomization helps.




randomized layouts without Zygotes



Zygotes defeat randomization


Issue 2: Threads

When a thread calls fork, only that thread goes. This can be problematic by itself, but is worse if locks are involved.




Issue 3: Incremental Init

Hierarchical Zygotess work when you can incrementally init more things mid-execution.

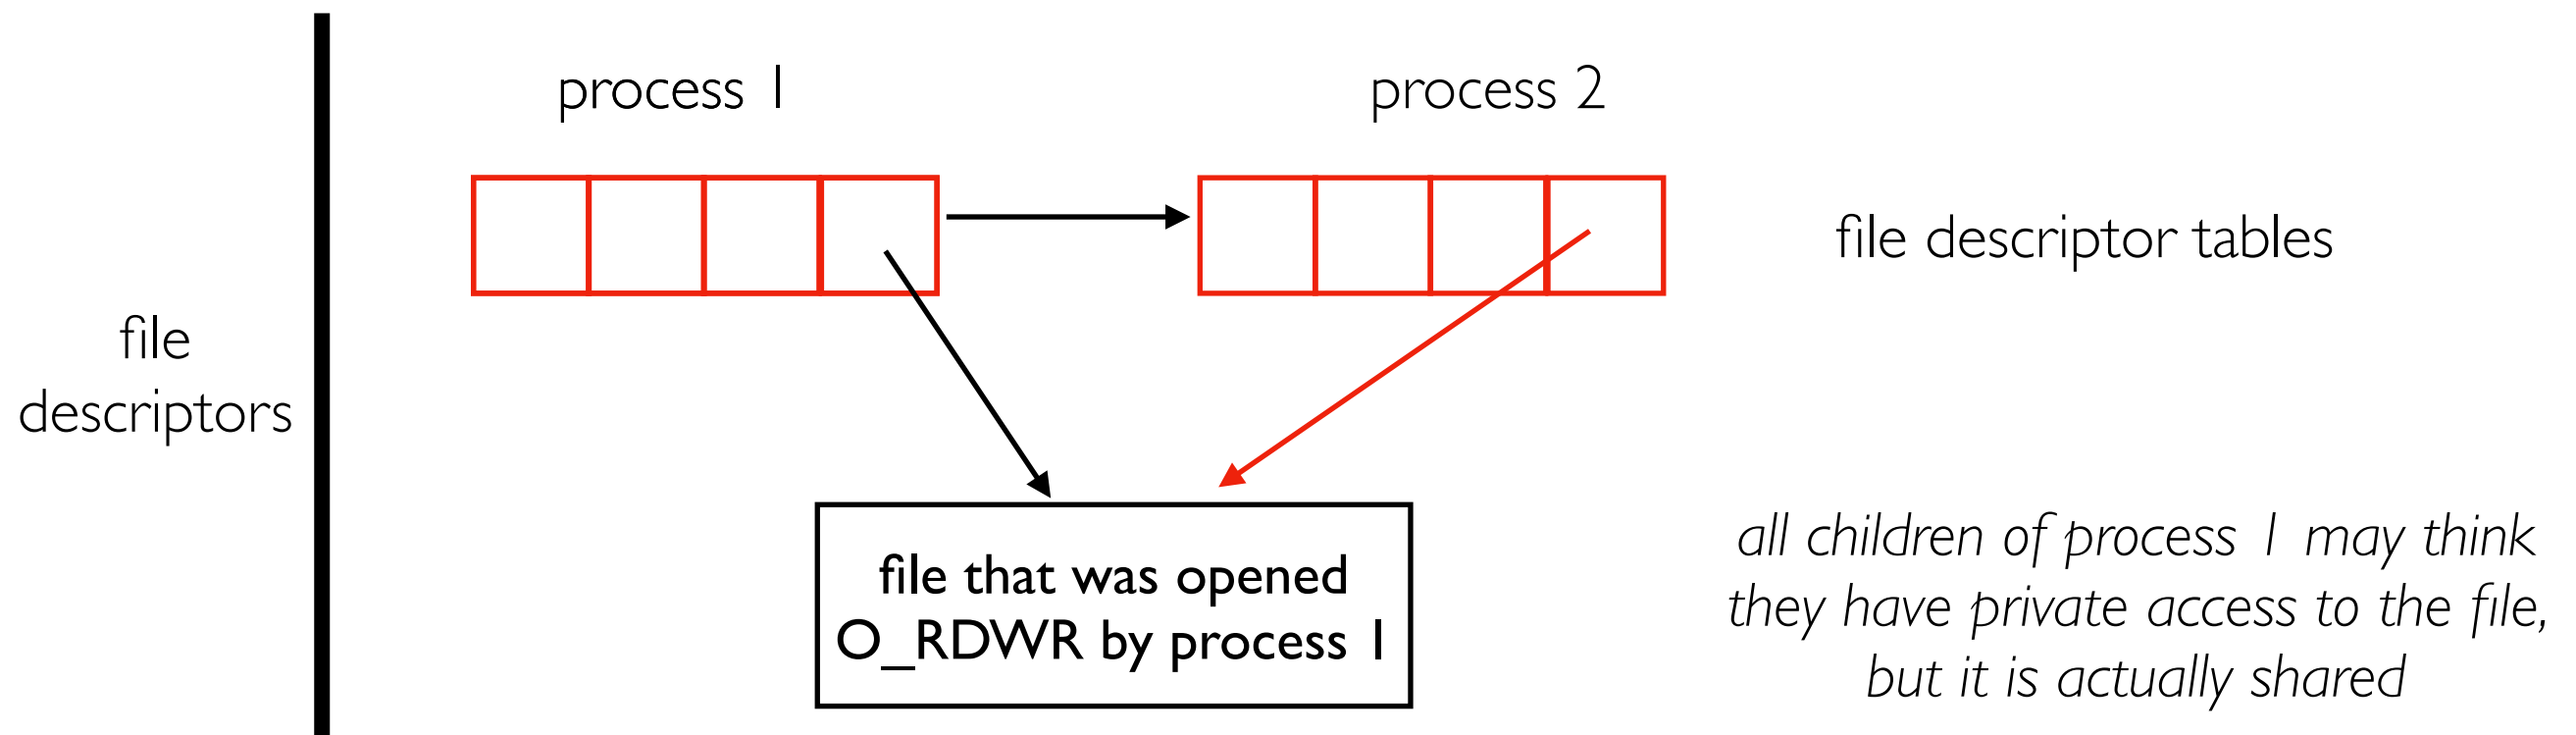


This wouldn't work if you're statically linking in all libraries (like in Rust or Go).


Issue 4: Unintentionally Shared Resources



fork has a similar effect on both virtual mem and FDs



Issue 5: Accounting



We're using 12 KB (3 pages) of physical memory total. How should we do accounting across the two processes?

- 6 KB each? If process 2 stops, does the memory accounting for process 1 pop, leading to a kill?
- 12 KB for process 1, since it allocated it first? Why should process 2 have a free ride?

Outline

Zygote Background

Zygote Limitations

Sandboxed Zygotes

OpenLambda Zygotes: Policy

Peeking Behind the Curtains of Serverless Platforms

Liang Wang, *UW-Madison*; Mengyuan Li and Yinqian Zhang, *The Ohio State University*;
Thomas Ristenpart, *Cornell Tech*; Michael Swift, *UW-Madison*

<https://www.usenix.org/conference/atc18/presentation/wang-liang>

1

AWS Lambda Instance Coldstart Latency >250ms

2

If N instances of a given lambda are idle,
AWS evicts $N/2$ instances every 5 minutes

Observation: we need to improve cold start latency
and/or make Lambda instances VERY memory efficient

SOCK: Rapid Task Provisioning with Serverless-Optimized Containers

Edward Oakes, Leon Yang, Dennis Zhou, and Kevin Houck, *University of Wisconsin-Madison*; Tyler Harter, *Microsoft, GSL*; Andrea C. Arpaci-Dusseau and Remzi H. Arpaci-Dusseau, *University of Wisconsin-Madison*

<https://www.usenix.org/conference/atc18/presentation/oakes>

Key Idea of OpenLambda/SOCK: start Lambdas by forking containerized **Zygotes** that have already imported certain dependencies (basic Zygote idea used by Android with JVM)

- 1 Improve startup time (pay import time only once)
- 2 Improve memory efficiency (copy-on-write sharing in kernel)

Parent Sandbox

Virtual



[python]

← cgroup memory **limit**

← memory **use**

Physical



very short background
on SOCK Zygotes

Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes

Parent Sandbox

Virtual



import pandas



Physical



Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes

Parent Sandbox

Virtual



import pandas

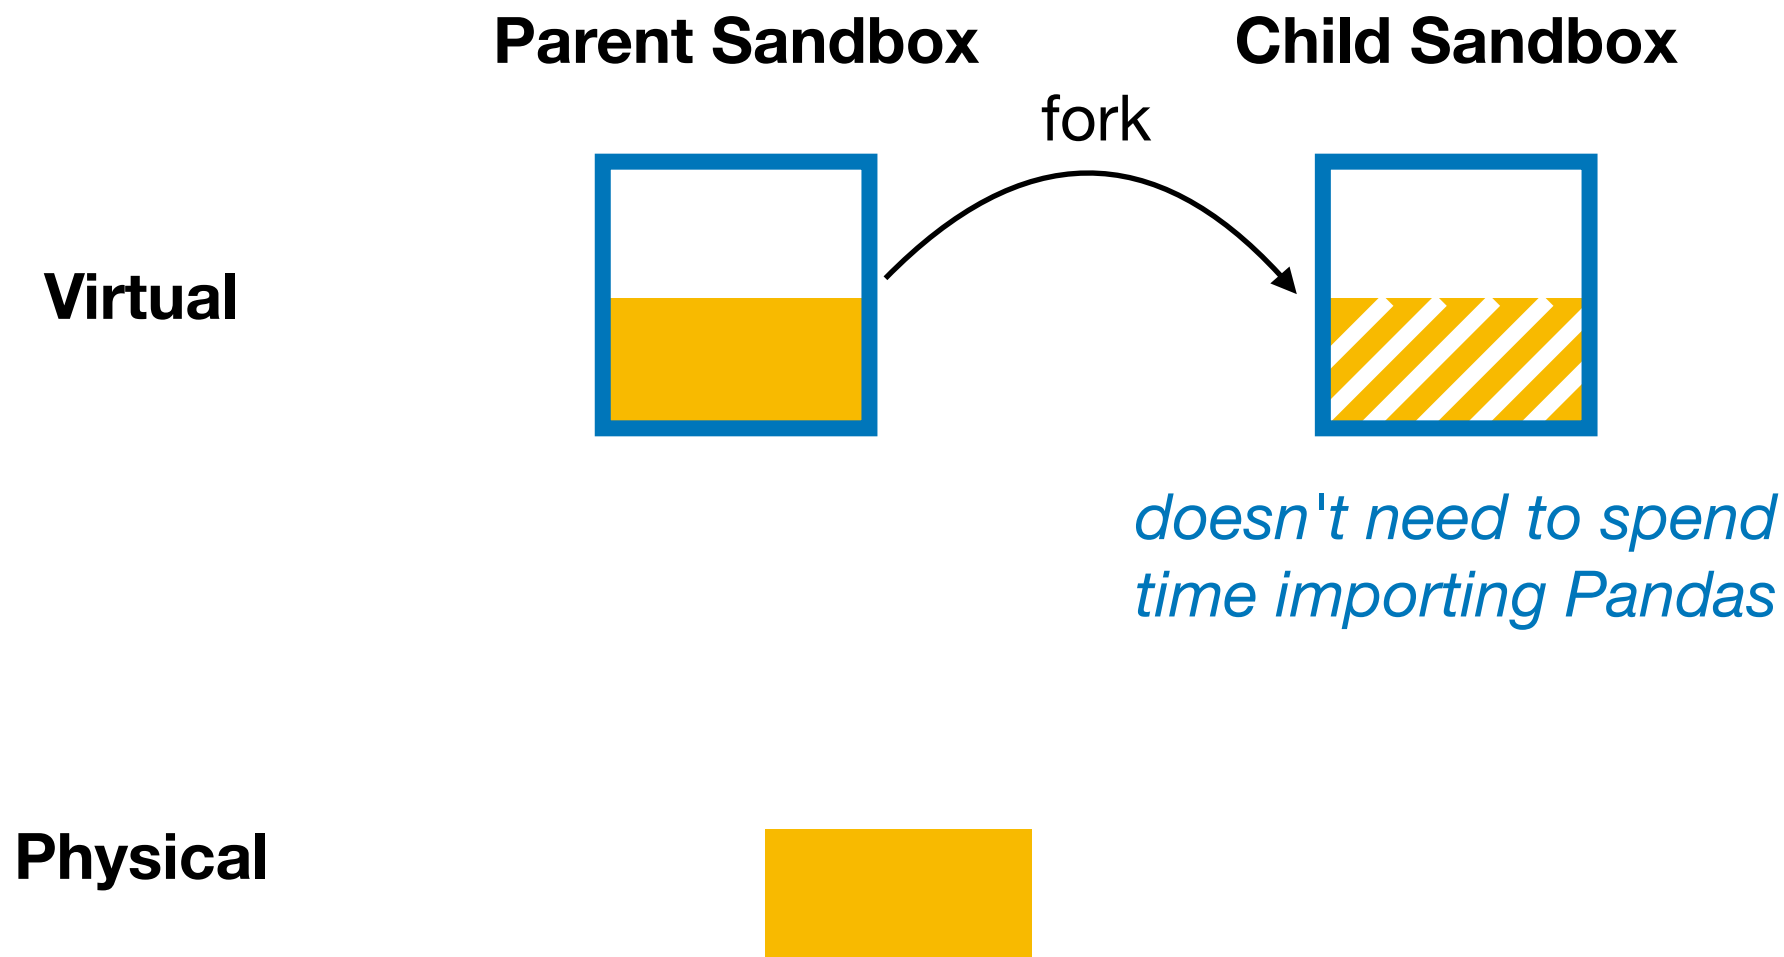


Physical



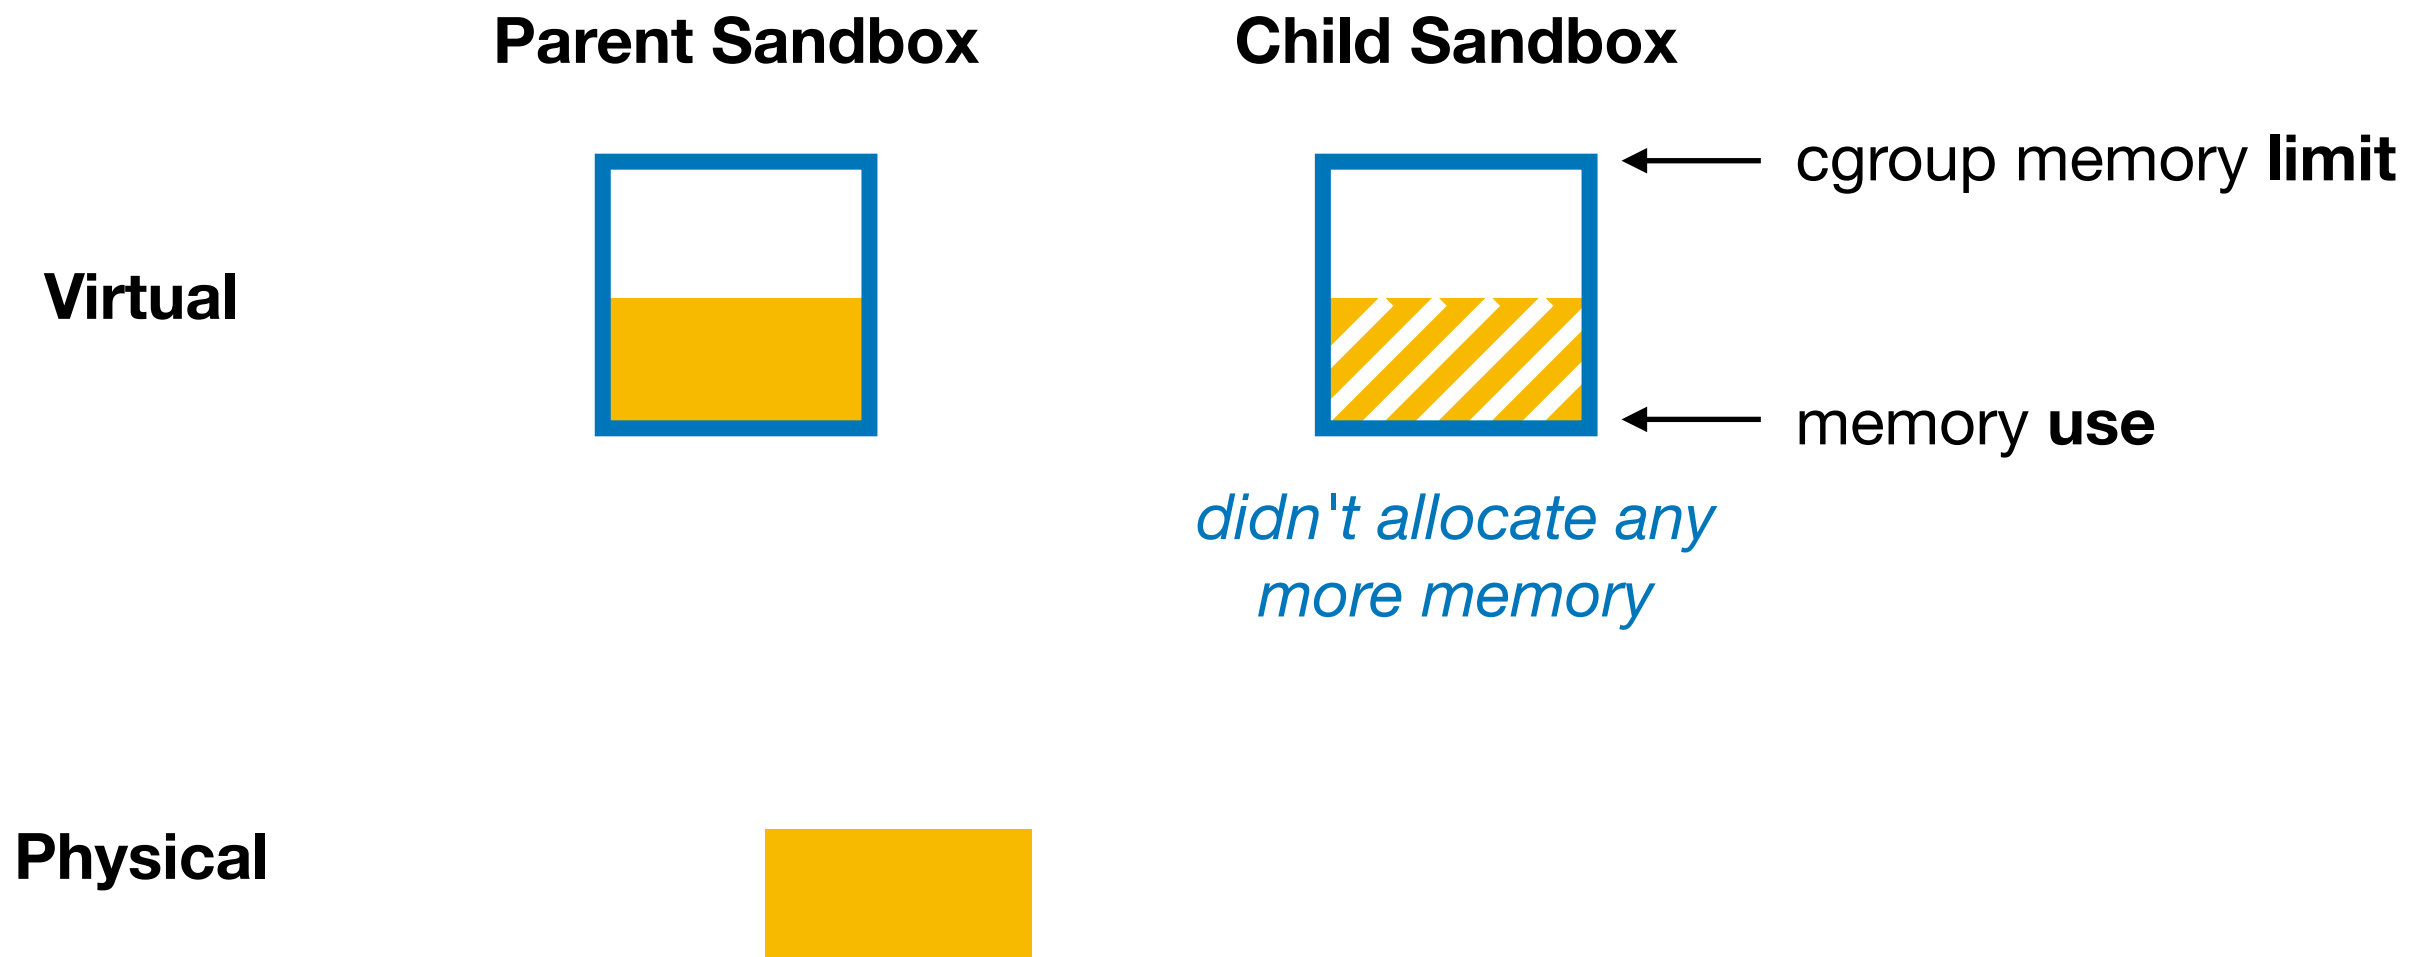
Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes




Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes




Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes



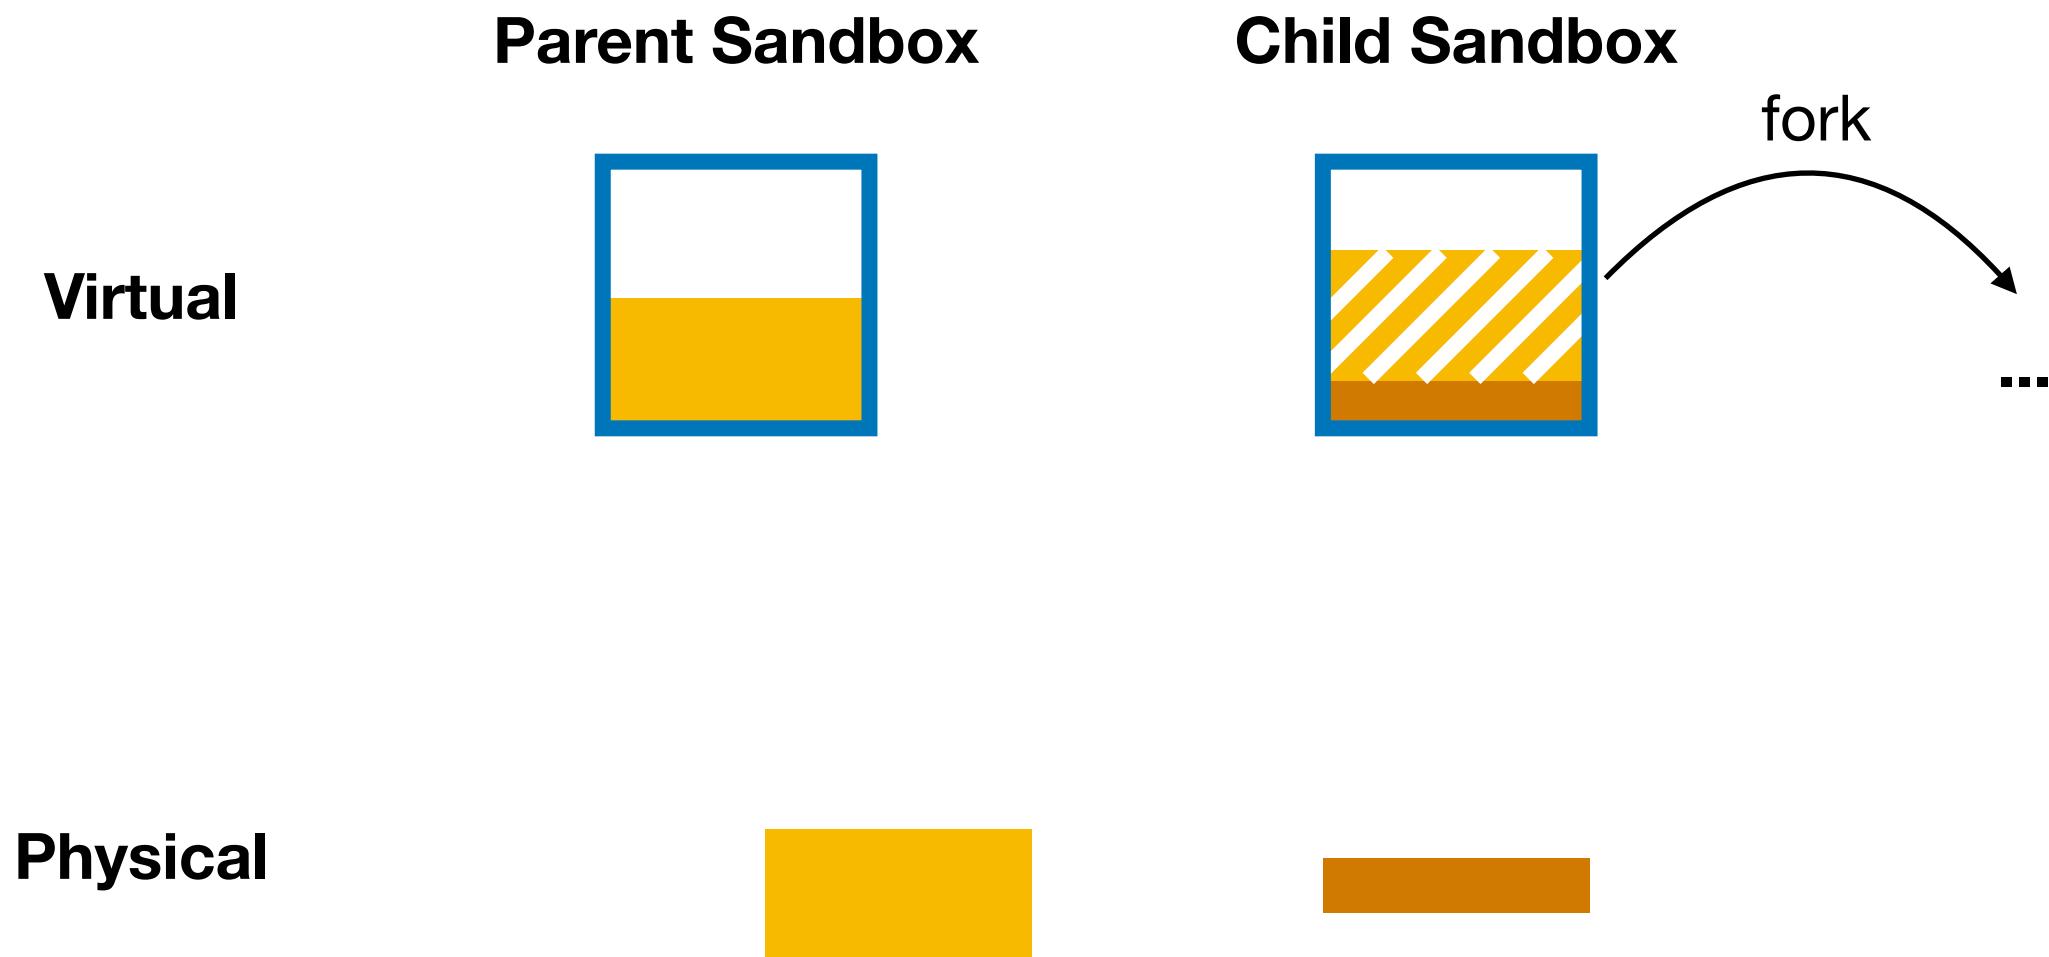
Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes



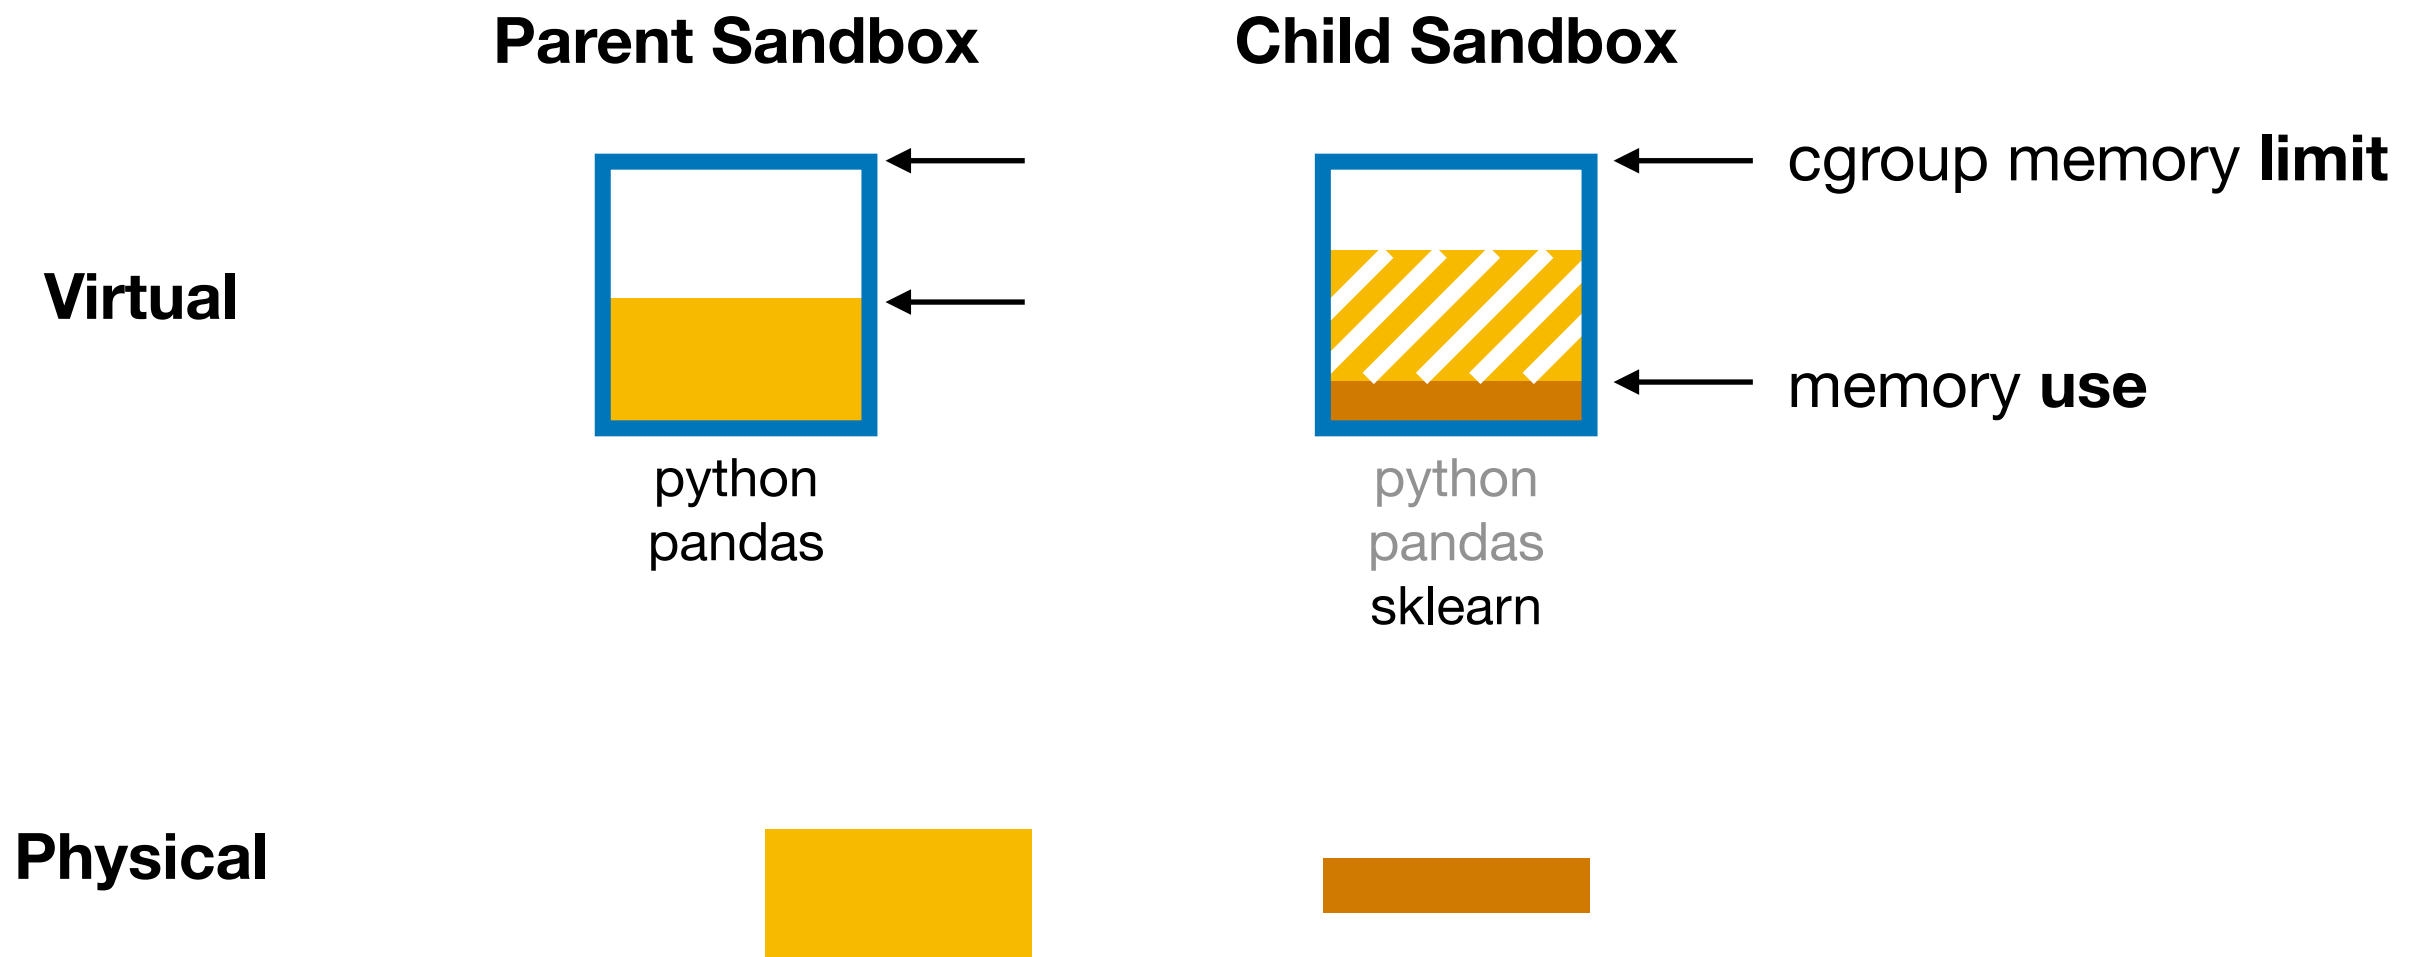
Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes



Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes



Mechanism: SOCK containers support sandbox forking

Policy: import cache maintains pool of "Zygote" containers with various pre-imports, providing parents for new sandboxes

Outline

Zygote Background

Zygote Limitations

Sandboxed Zygotes

OpenLambda Zygotes: Policy

Previous Cache Policy

Policy in SOCK paper was **dynamic**, but **shortsighted**:

- If a Zygote has exactly the packages needed by lambda, use it
- Else:
 1. find a Zygote **ZA** with "greatest" subset of packages
 2. create a Zygote **ZB** from **ZA**, adding any lacking packages
 3. use **ZB** to create Sandbox for lambda

Problem: cannot take advantage of certain patterns obvious across many different lambdas. **Pathological case:**

- there are N packages, and $N*(N-1)/2$ lambdas
- each lambda imports a different pair of packages
- a new Zygote will be created each time, but it will never be used more than once!

Inspiration for a better Zygote tree policy

1

Cache-aware load balancing of data center applications

Aaron Archer
Google
New York, New York
aarcher@google.com

Vahab Mirrokni
Google
New York, New York
mirrokni@google.com

Kevin Aydin
Google
New York, New York
kaydin@google.com

Aaron Schild
UC Berkeley
Berkeley, California
aschild@berkeley.edu

MohammadHossein Bateni
Google
New York, New York
bateni@google.com

Ray Yang
Google
New York, New York
rayy@google.com


Idea: statically compute the cache structure offline (e.g., from last night's traces); tweak to make it somewhat dynamic

2

Decision Trees

Idea: greedily build a tree by starting with one node, specifying a metric for the quality of the tree, and greedily split nodes to improve quality the most

Inspiration for a better Zygote tree policy




compute gray structure offline, create/evict black Zygotes based on workload

2 Decision Trees

Idea: greedily build a tree by starting with one node, specifying a metric for the quality of the tree, and greedily split nodes to improve quality the most

Inspiration for a better Zygote tree policy



compute gray structure offline, create/evict black Zygotes based on workload

2 Decision Trees

Idea: greedily build a tree by starting with one node, specifying a metric for the quality of the tree, and greedily split nodes to improve quality the most

Formulating Lambda Trace for Decision Tree

Matrix constructed
from last night's trace

Root Node					
	A	B	C	D	Y
fn1	1	1	0	0	1
fn2	1	0	1	0	1
fn3	1	0	0	1	1
fn4	0	1	0	1	0
fn5	0	0	1	1	1
fn6	0	0	0	1	0
fn7	0	1	0	0	0
fn8	0	0	0	0	0

don't overthink what the "Y"
is that we're trying to predict
(it will go away soon)

Lambda function **fn5** imports packages C and D

Formulating Lambda Trace for Decision Tree

Root Node

	A	B	C	D	Y
fn1	1	1	0	0	1
fn2	1	0	1	0	1
fn3	1	0	0	1	1
fn4	0	1	0	1	0
fn5	0	0	1	1	1
fn6	0	0	0	1	0
fn7	0	1	0	0	0
fn8	0	0	0	0	0

This one-node tree has a high *impurity score*(*), because there is a 50/50 mix of $Y=0$ and $Y=1$ in the same node (the worst case)

*common impurity metrics: Gini, entropy, variance

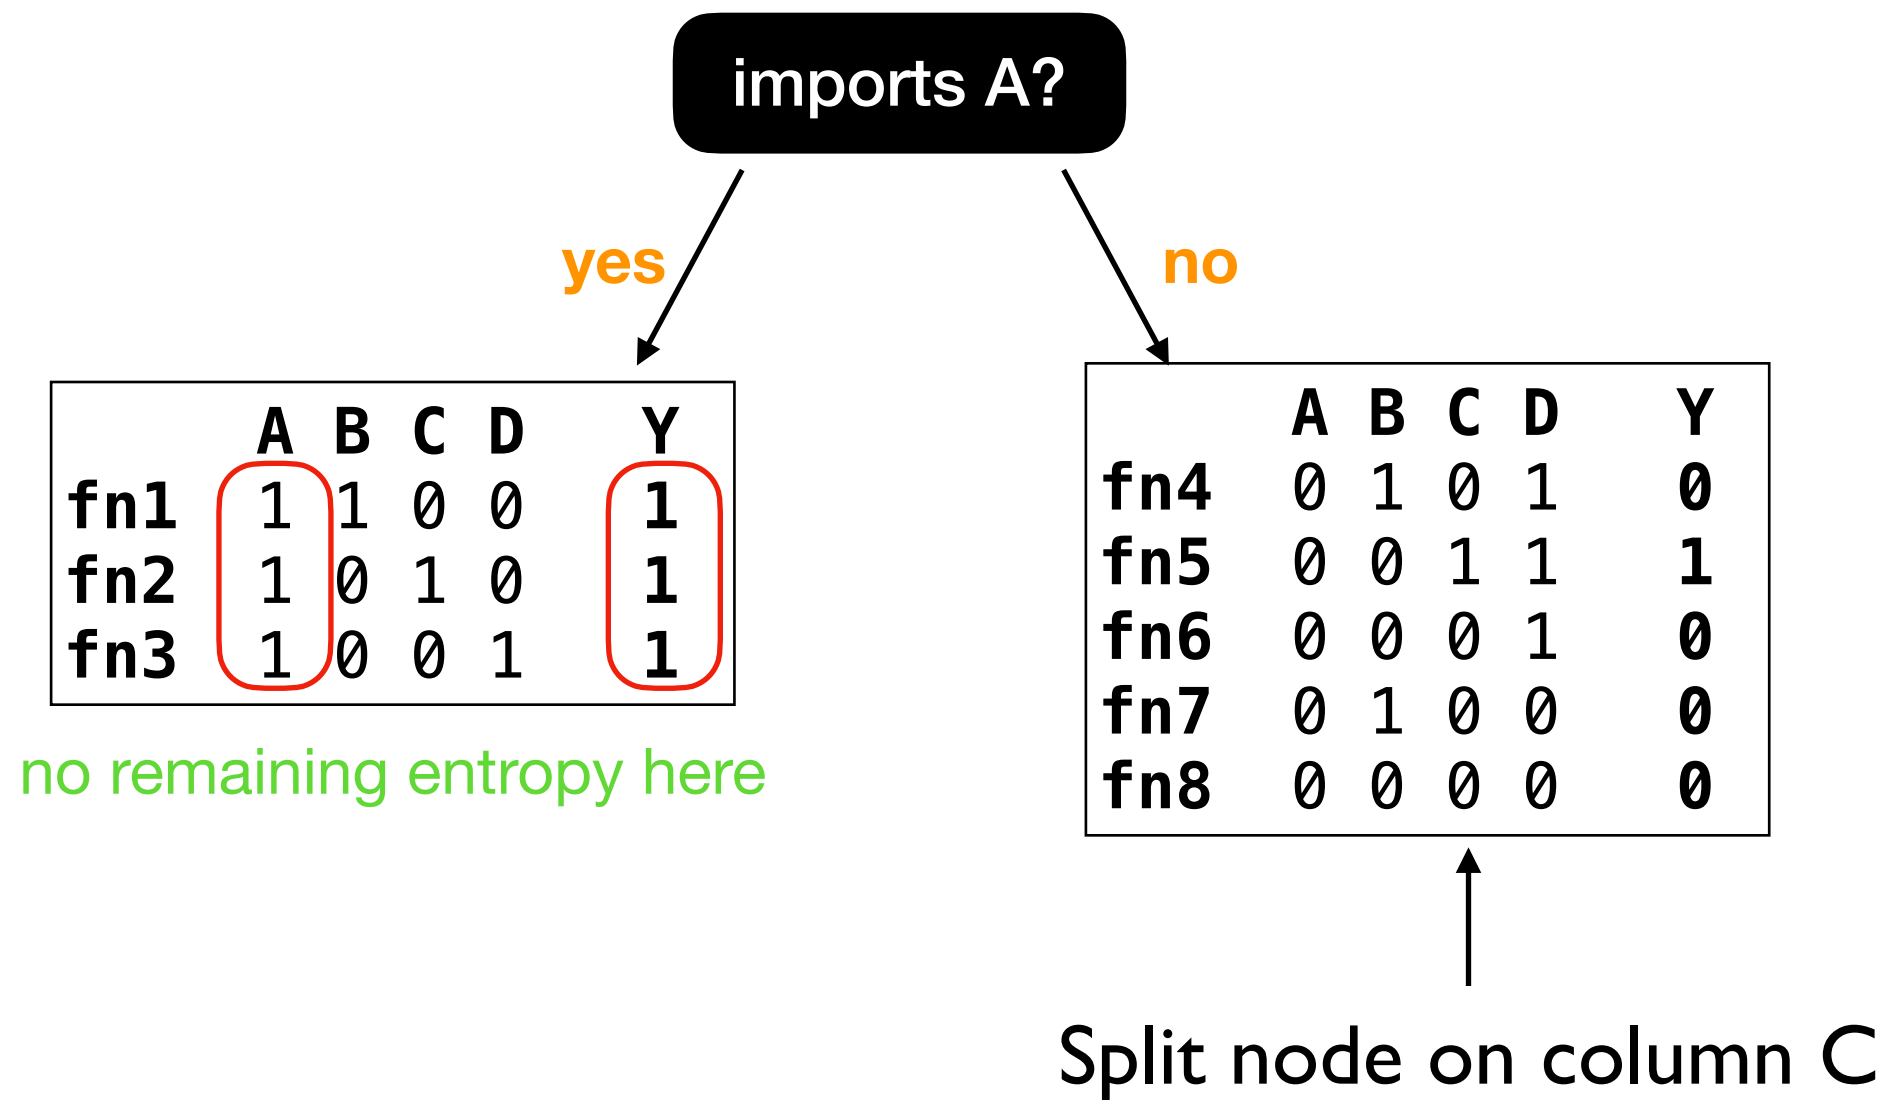
Formulating Lambda Trace for Decision Tree

Root Node


	A	B	C	D	Y
fn1	1	1	0	0	1
fn2	1	0	1	0	1
fn3	1	0	0	1	1
fn4	0	1	0	1	0
fn5	0	0	1	1	1
fn6	0	0	0	1	0
fn7	0	1	0	0	0
fn8	0	0	0	0	0

↑
Split node on column A in root node,
because that will organize Y the most

Formulating Lambda Trace for Decision Tree

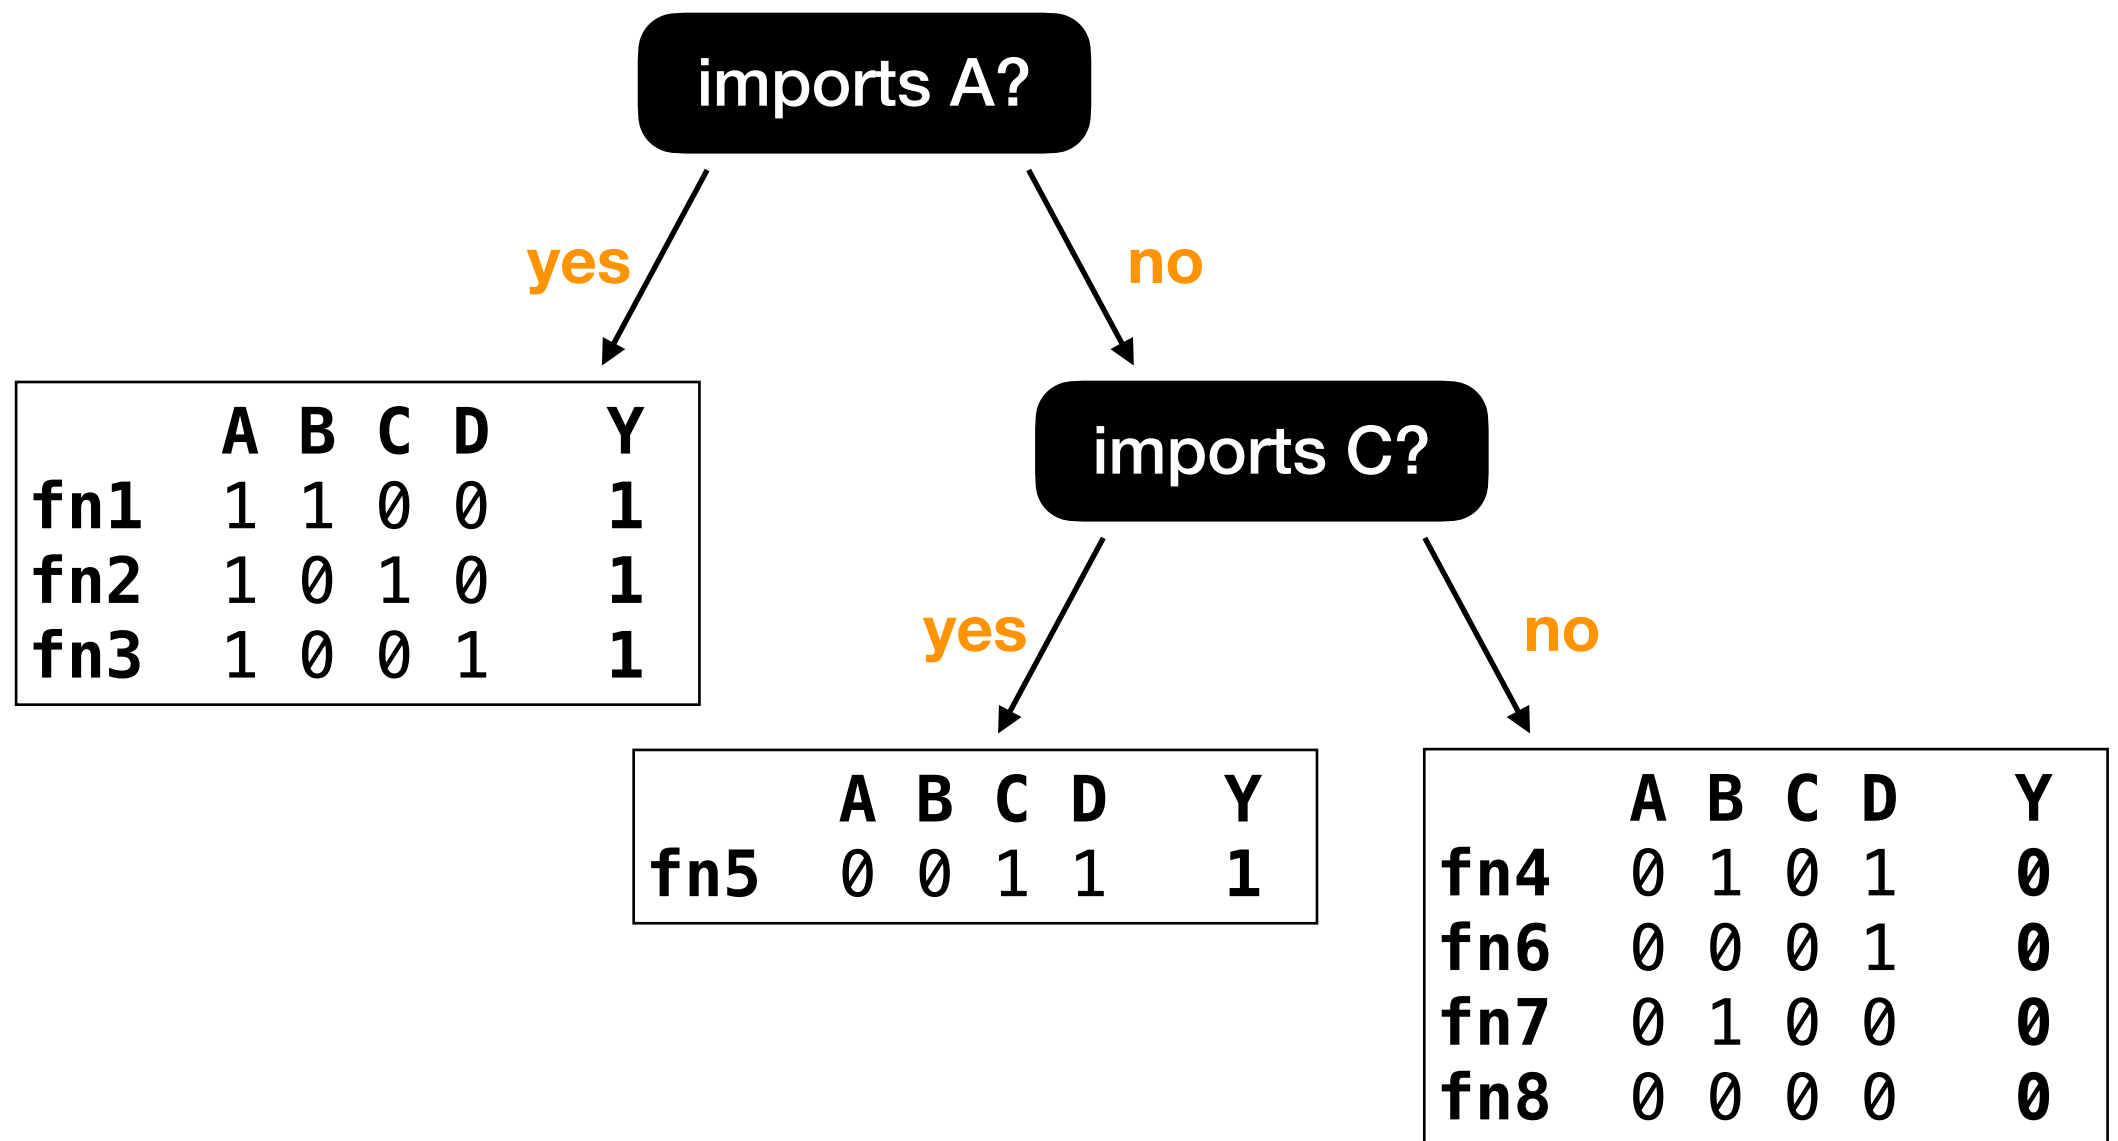


Formulating Lambda Trace for Decision Tree



This perfectly fits the training data!

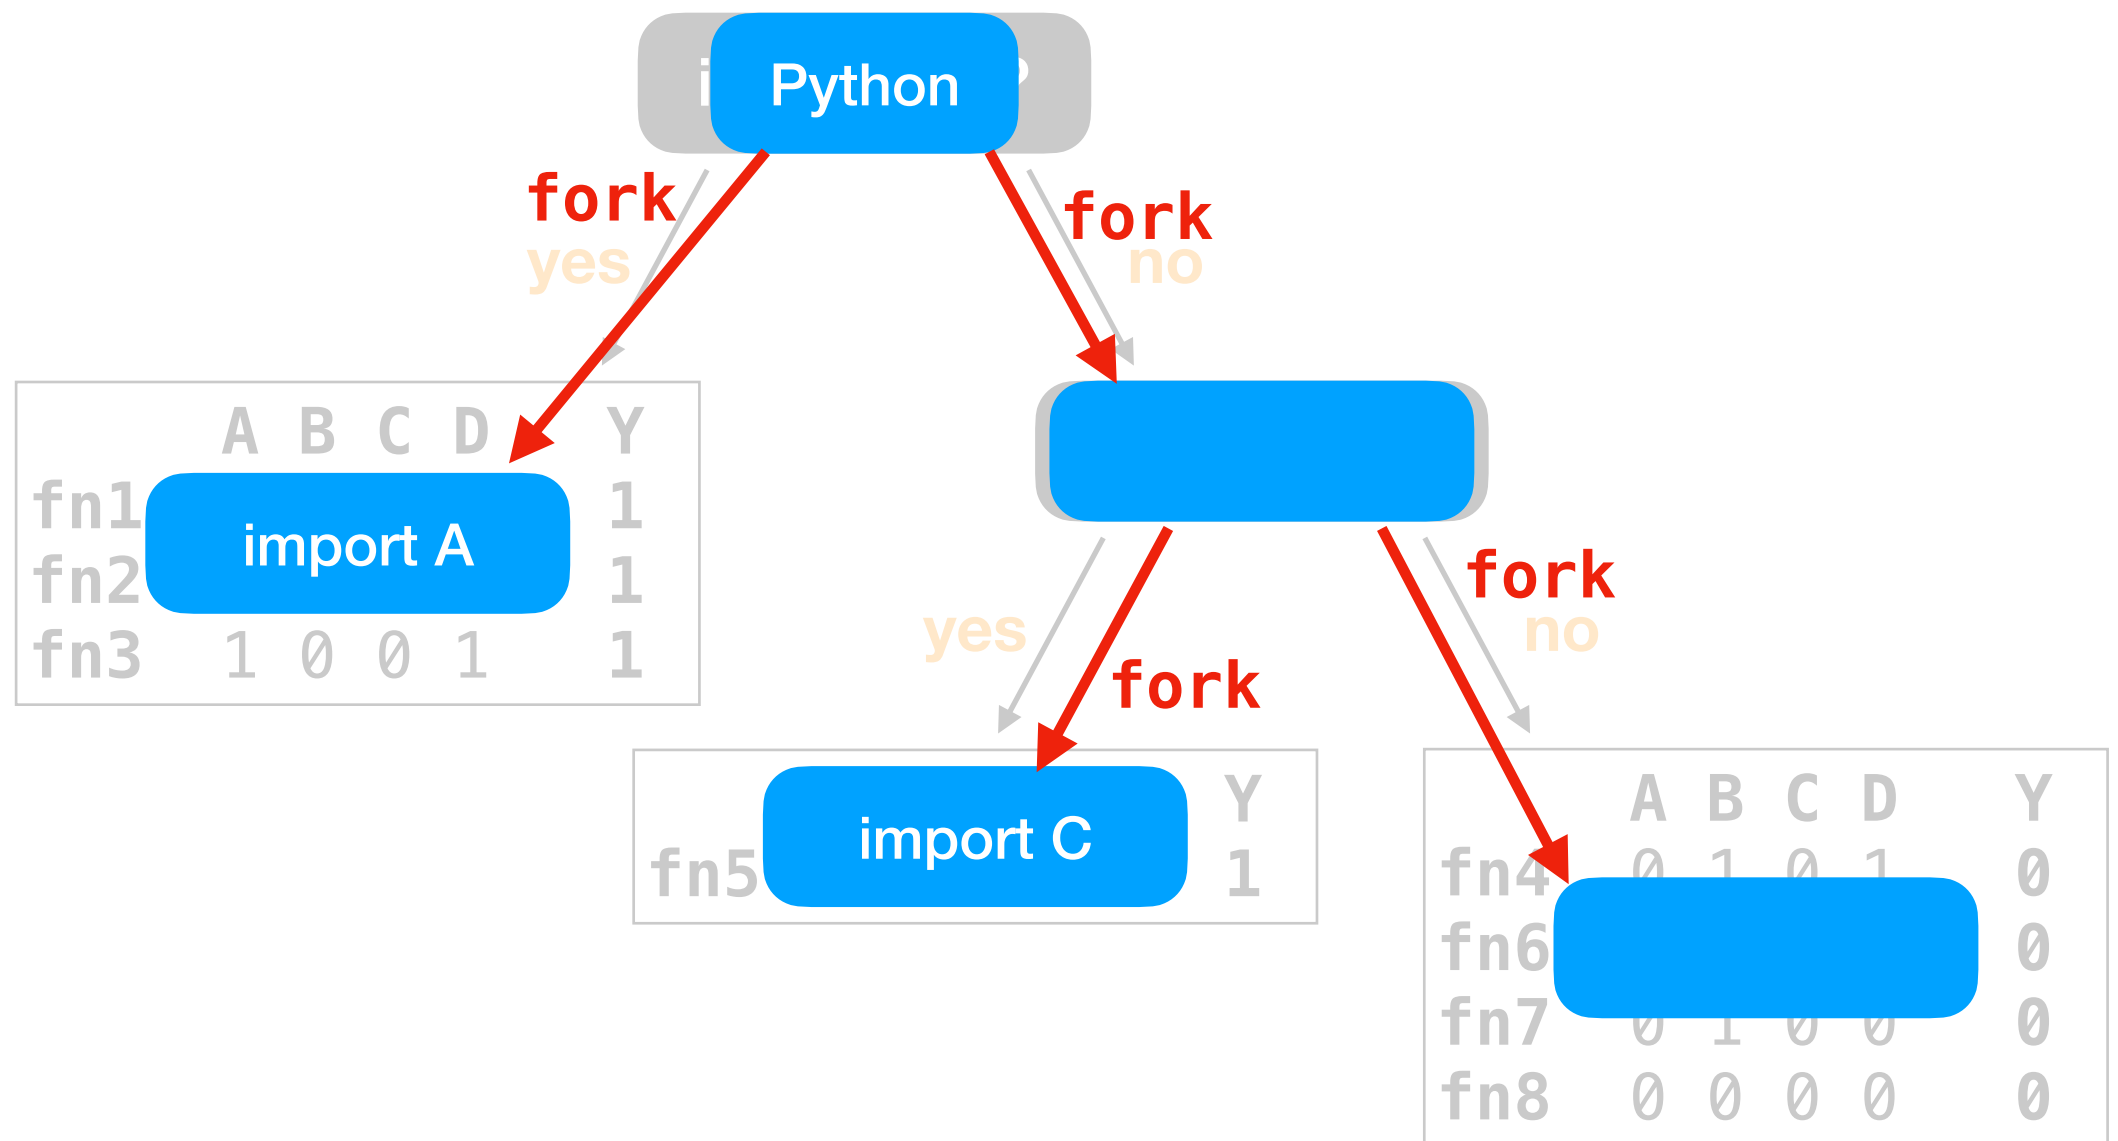
Formulating Lambda Trace for Decision Tree



This perfectly fits the training data!

Note: we *could* also construct a **5-node Zygote tree** from this decision tree

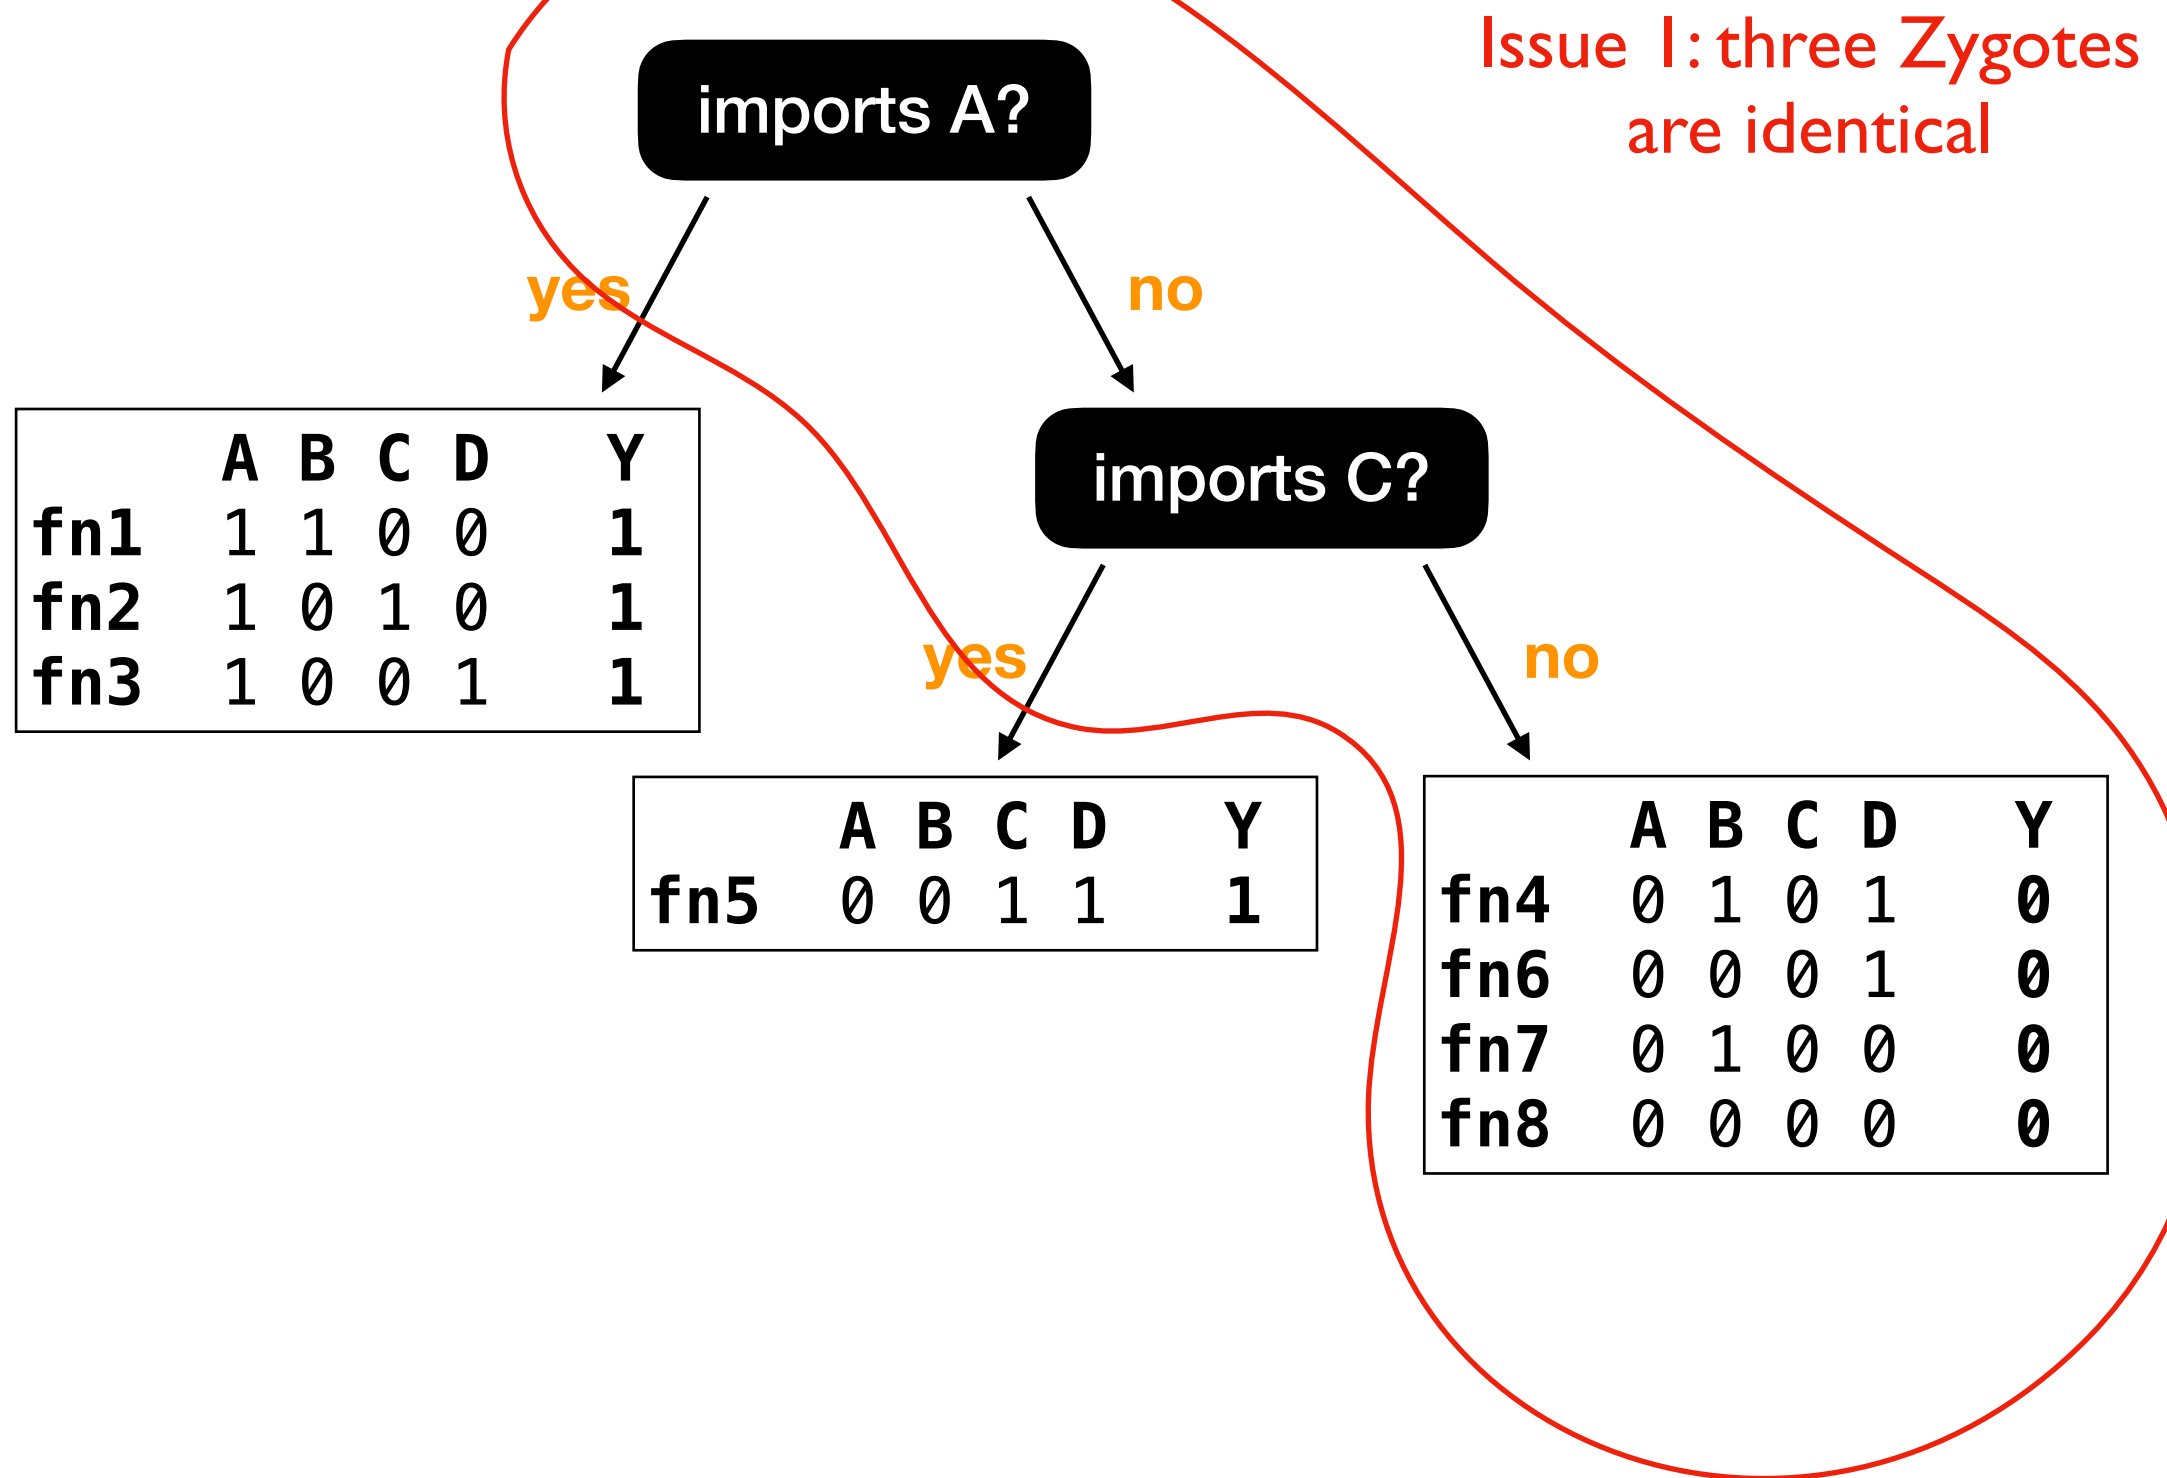
Formulating Lambda Trace for Decision Tree



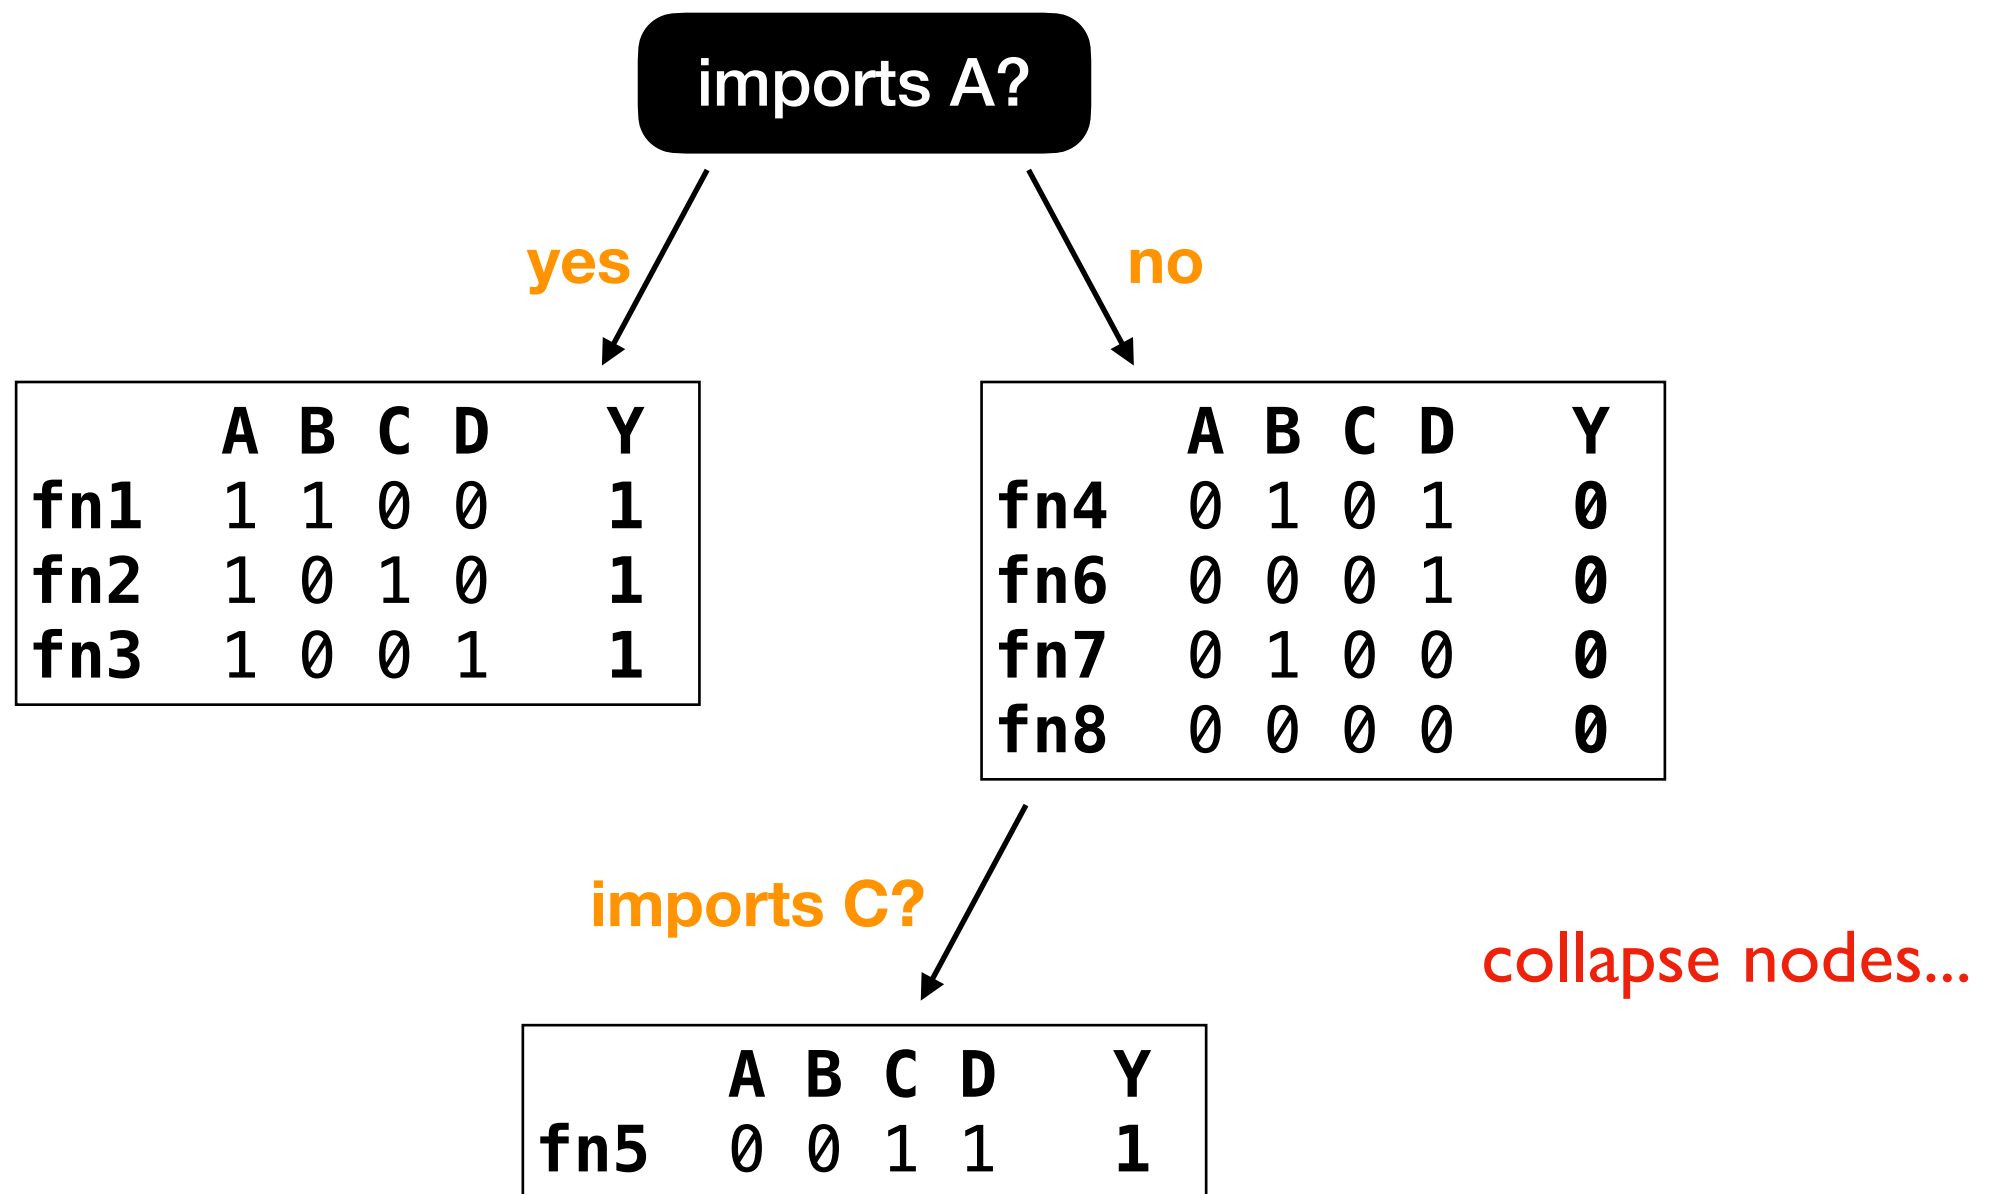
This perfectly fits the training data!

Note: we *could* also construct a 5-node Zygote tree from this decision tree

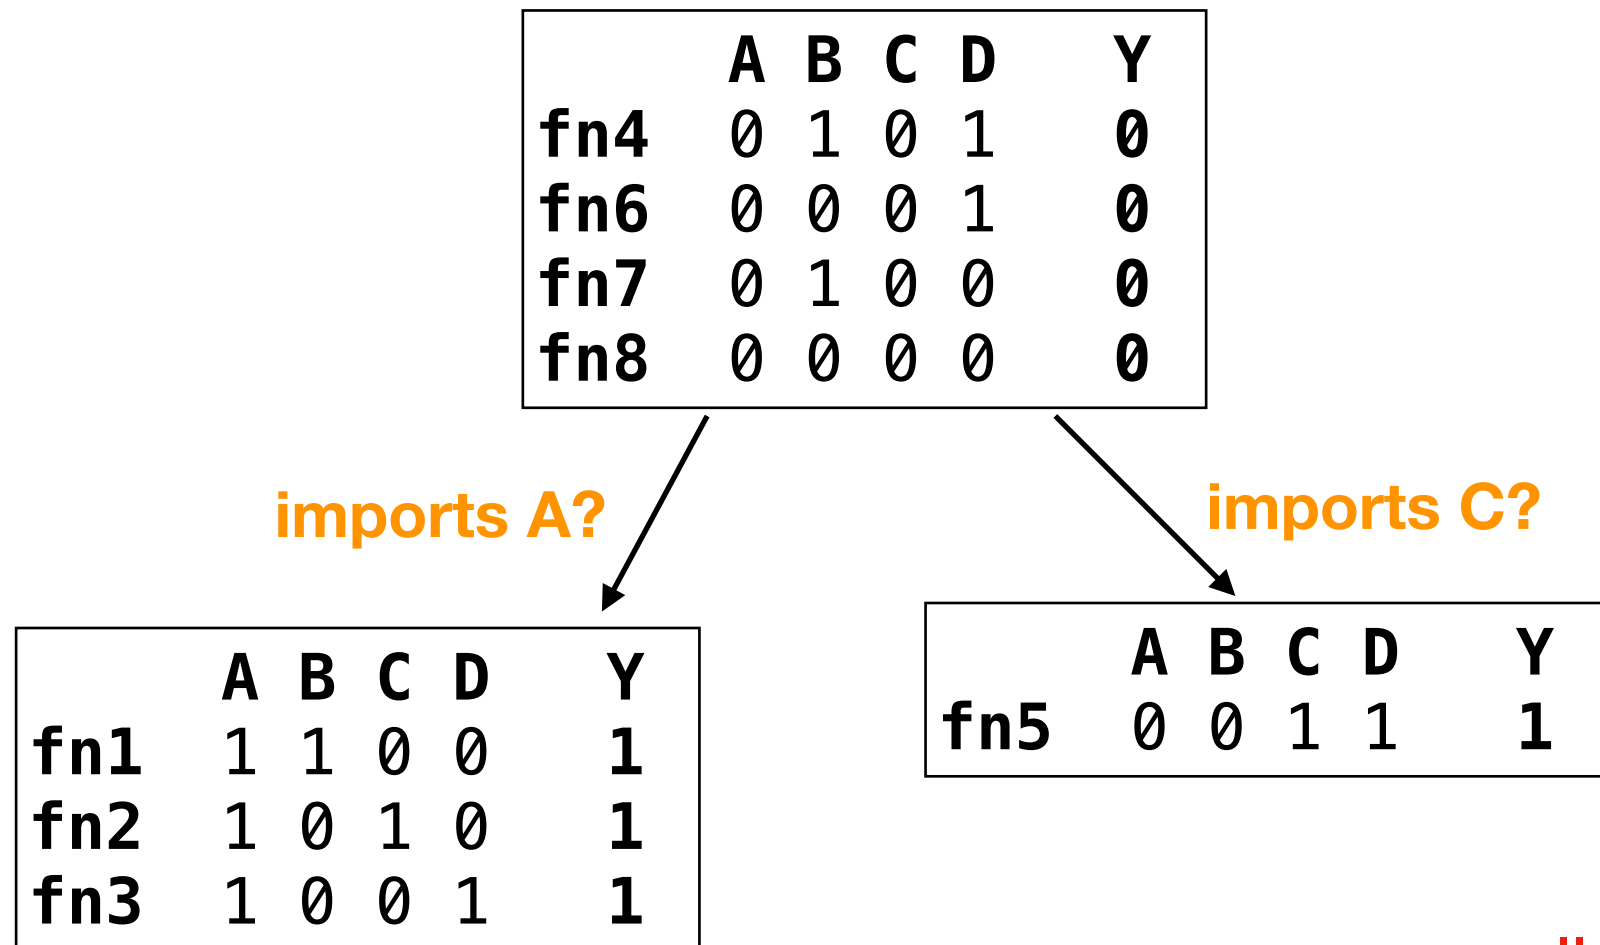
Formulating Lambda Trace for Decision Tree



Formulating Lambda Trace for Decision Tree

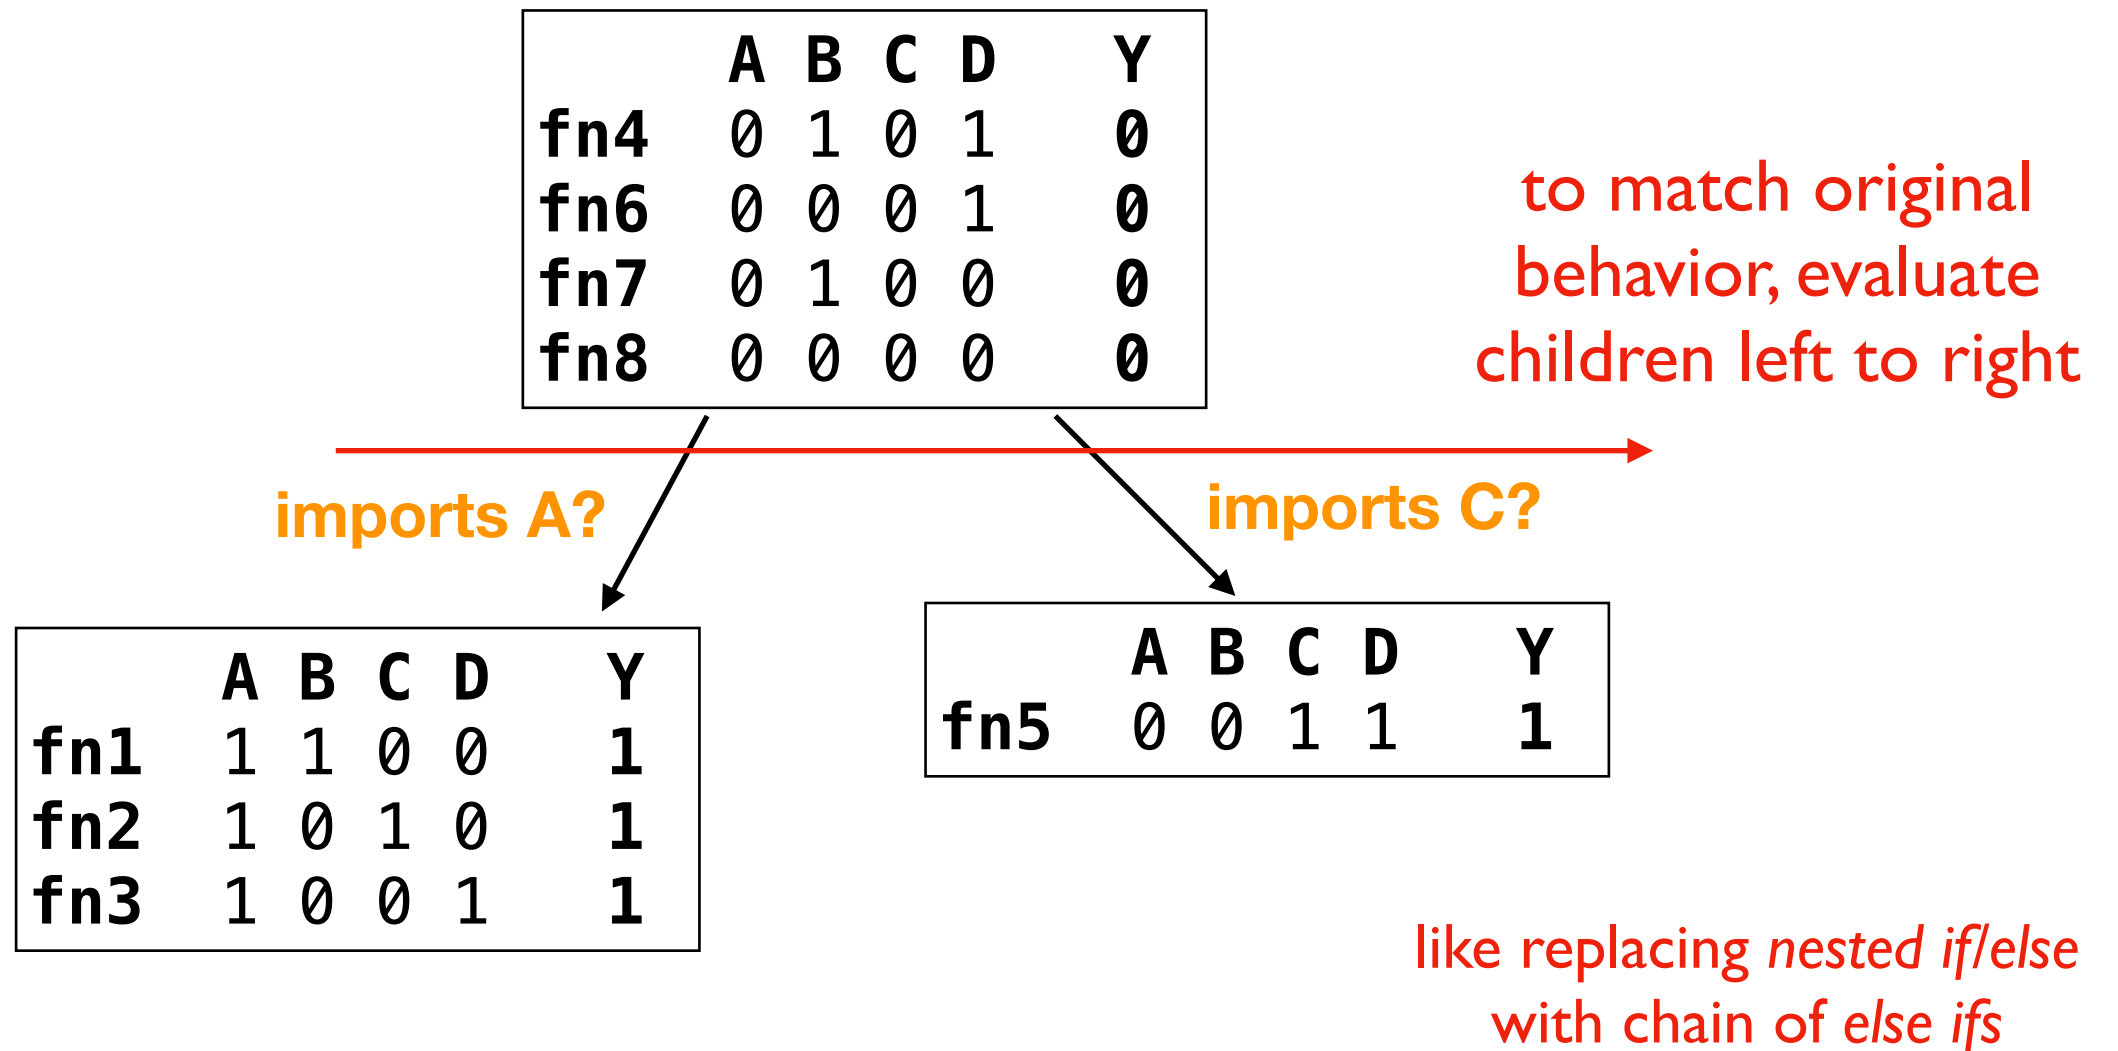


Formulating Lambda Trace for Decision Tree

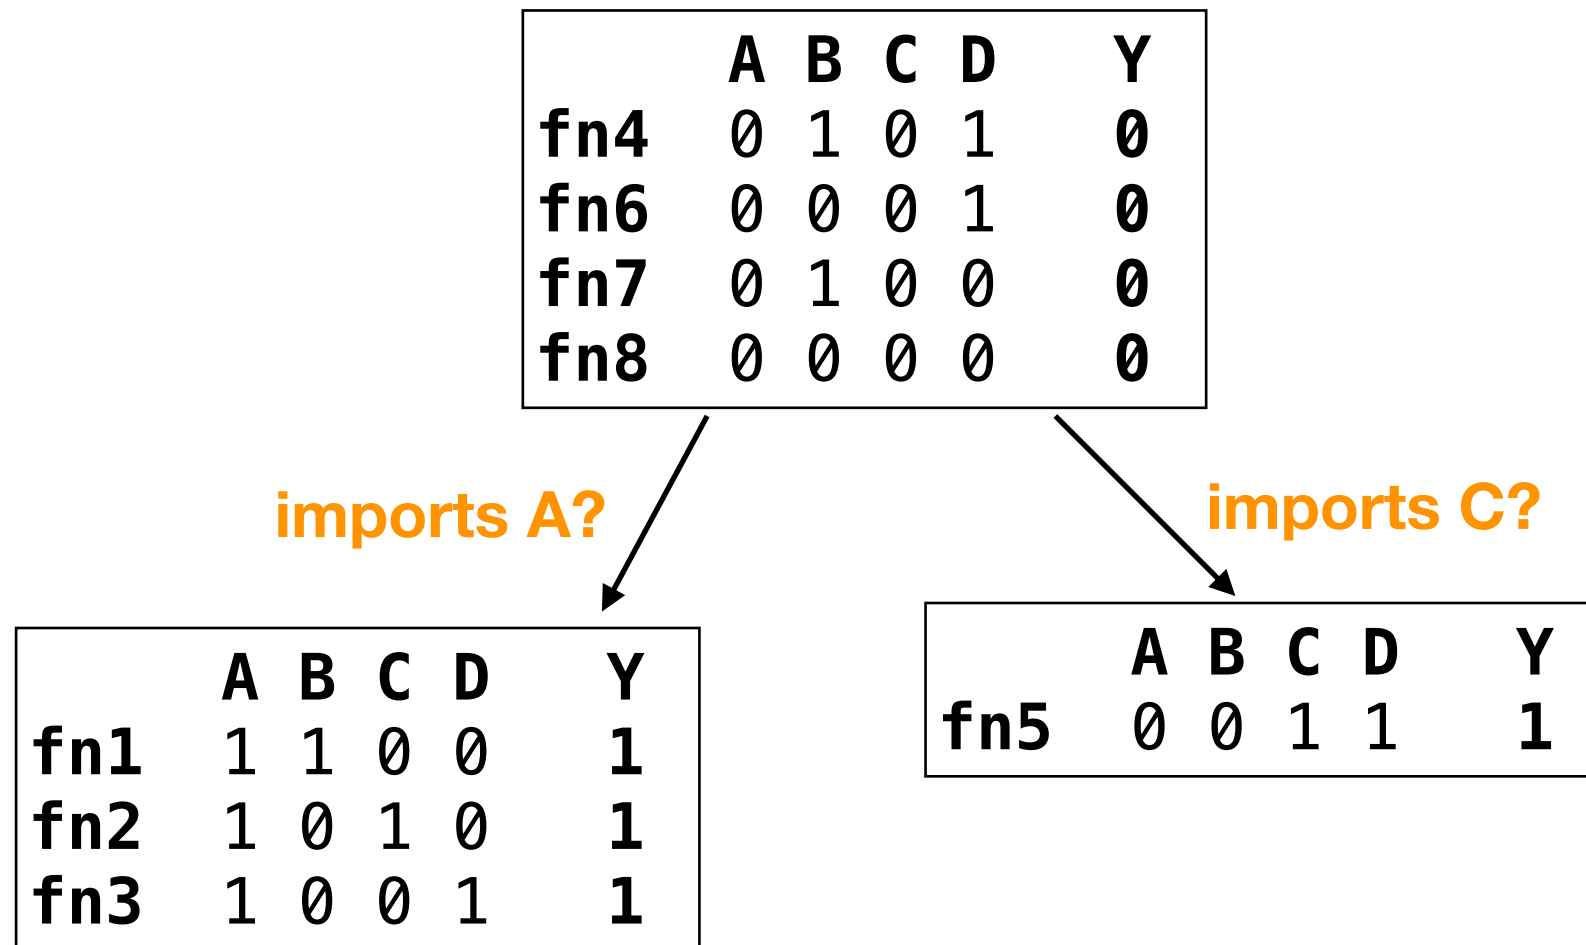


collapse nodes...

Formulating Lambda Trace for Decision Tree

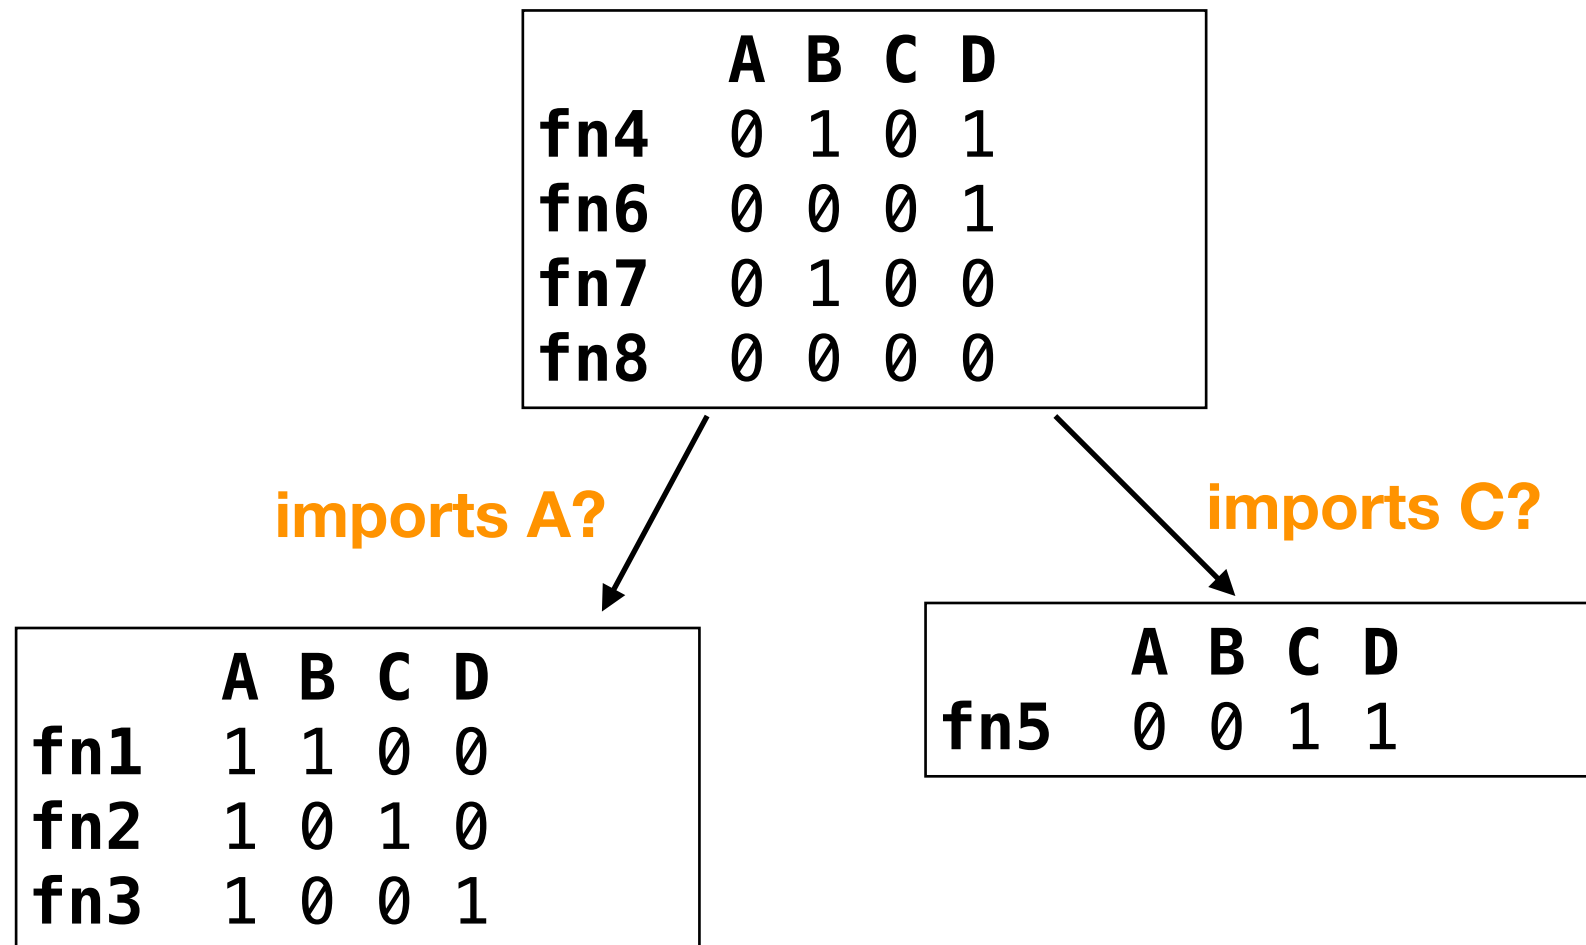


Formulating Lambda Trace for Decision Tree



Issue 2: what is Y anyway?

Formulating Lambda Trace for Decision Tree




Issue 2: what is Y anyway?

We don't need it if we invent our own perf-oriented impurity measure.
We're not trying to classify anything!


we'll propose 5 components for such a metric

Rule 0: Prioritize common packages



	pkg A	pkg B
fn 1	1	1
fn 2	1	0
fn 3	1	0
fn 4	0	0

Rule 0: Prioritize common packages



	<i>pkg A</i>	<i>pkg B</i>	"Y"
fn 1	1	1	1
fn 2	1	0	0
fn 3	1	0	0
fn 4	0	0	1

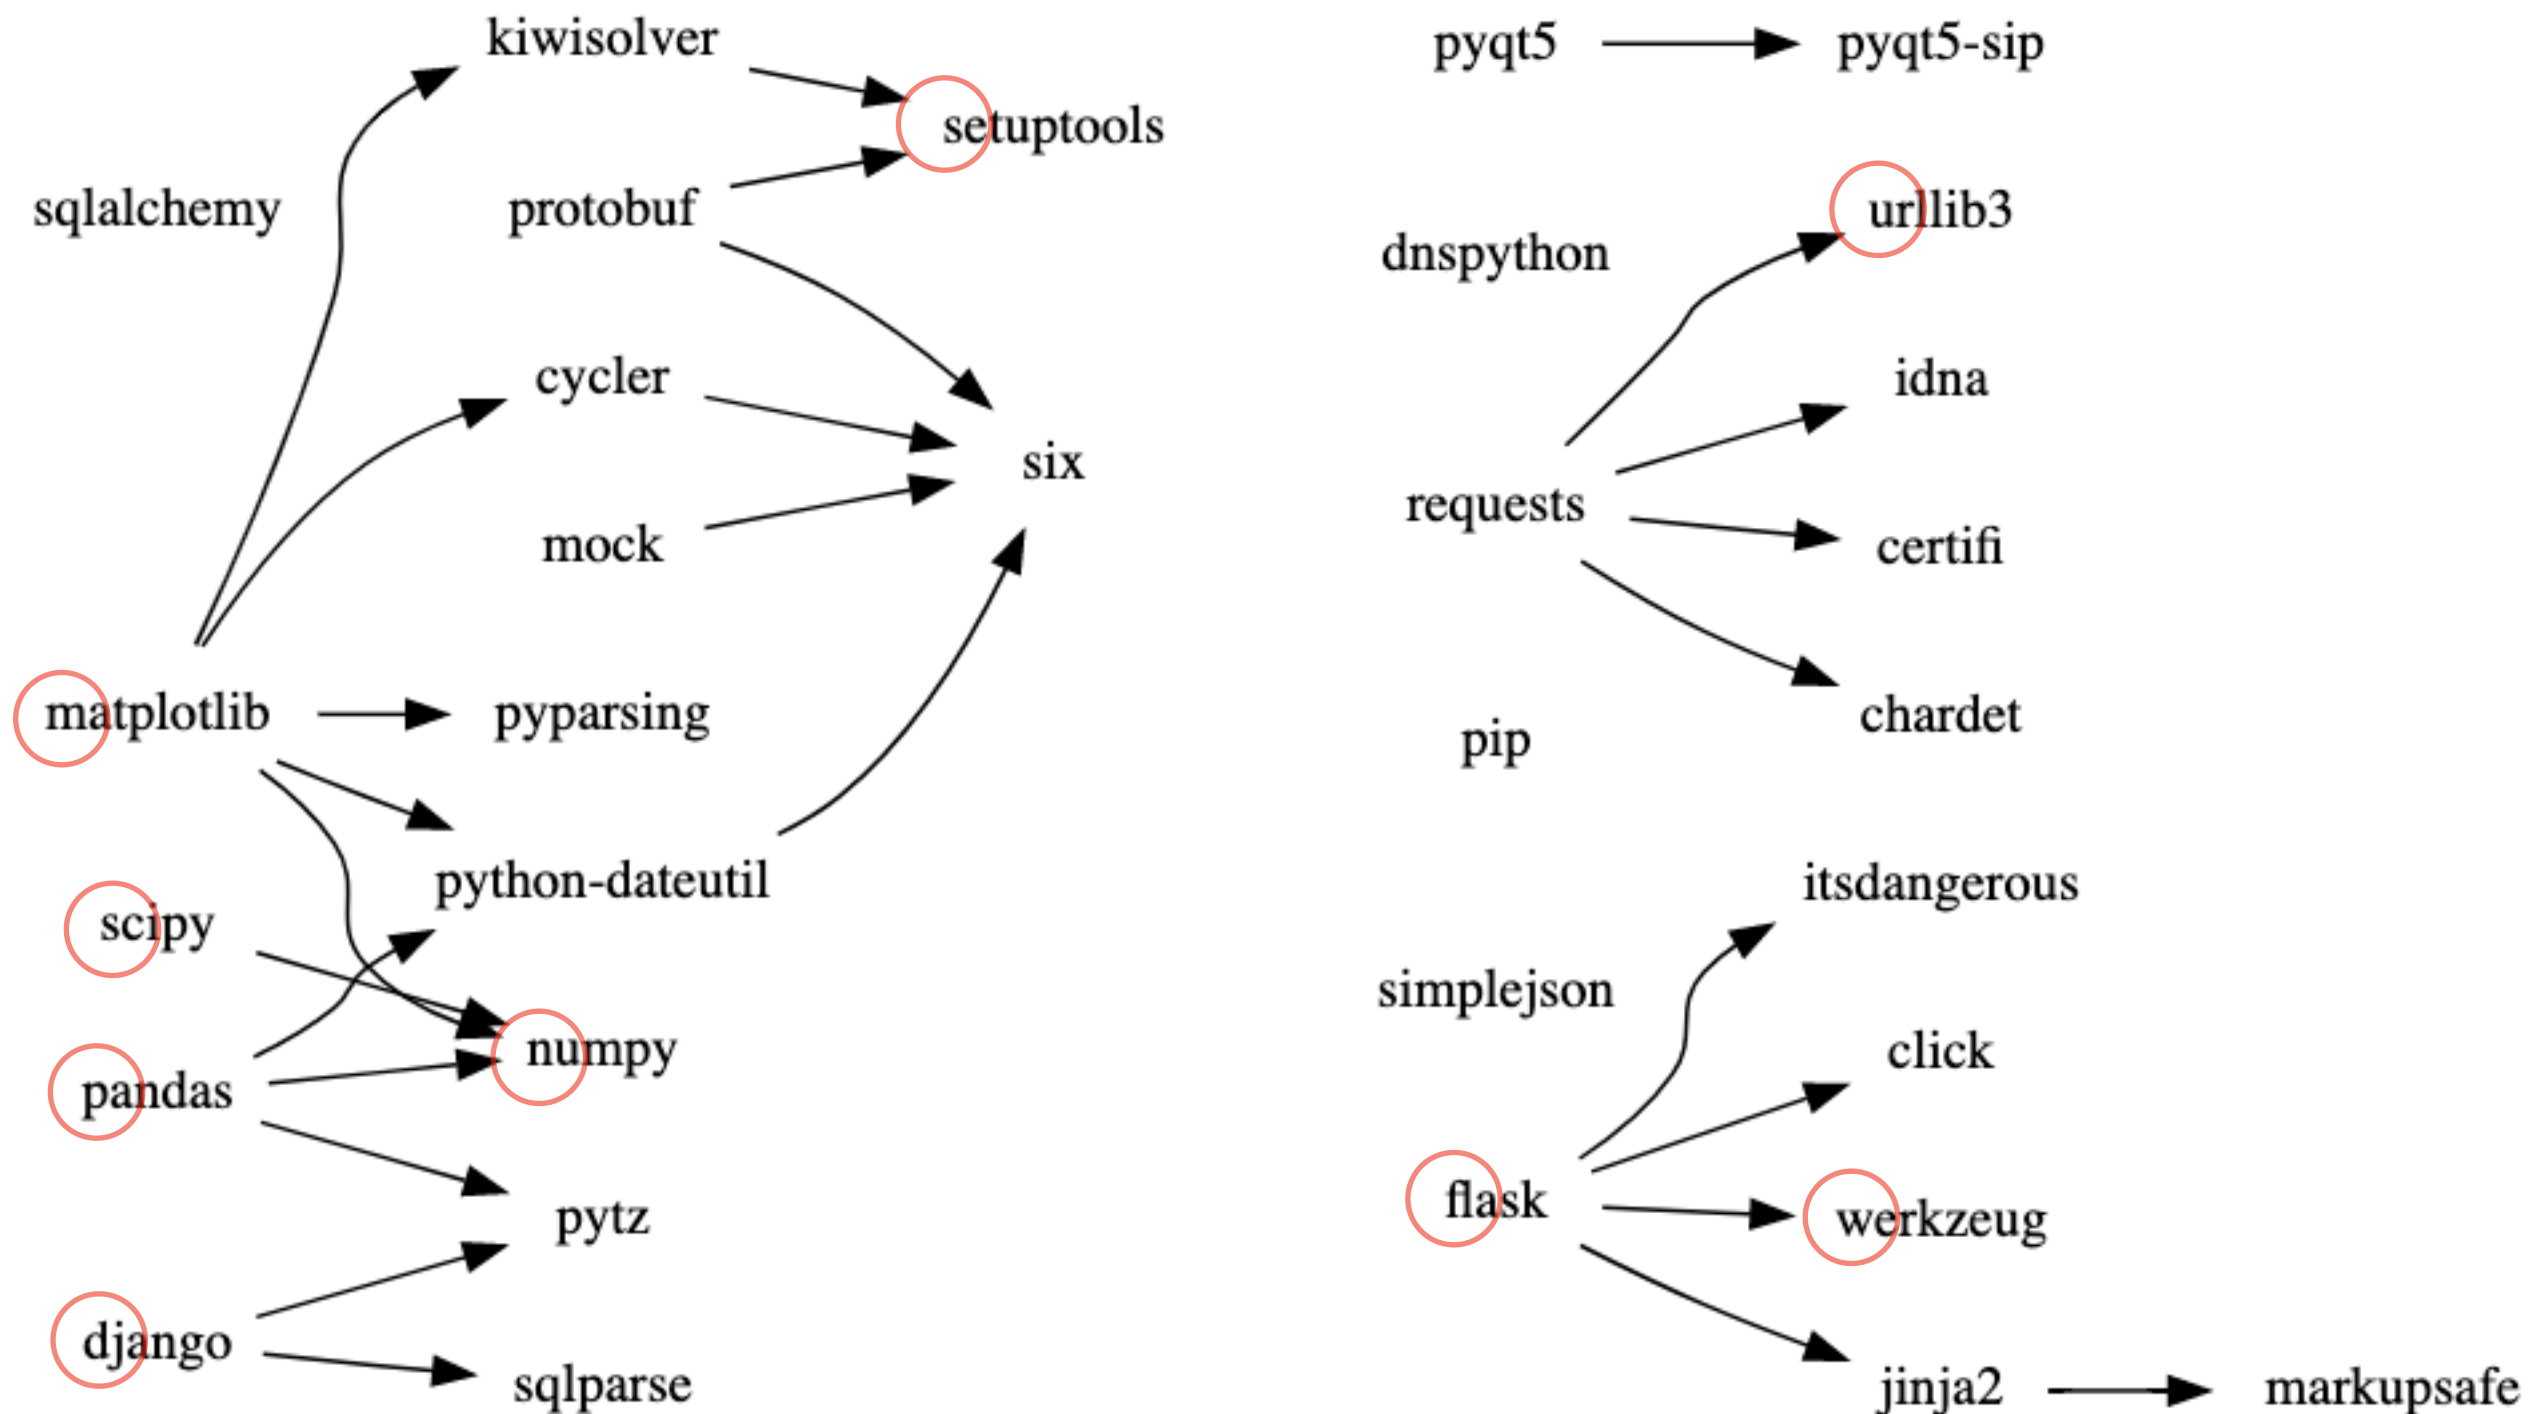
if we *were* classifying, note that we wouldn't have any bias for/against 1 or 0

Evaluation: Workload Generator

32 Real Packages

- Start with 20 top packages (SOCK paper)
- Eliminate 2 that don't work with OpenLambda currently
- Add 14 dependencies

Benchmark Packages and Dependencies



Evaluation: Workload Generator

32 Real Packages

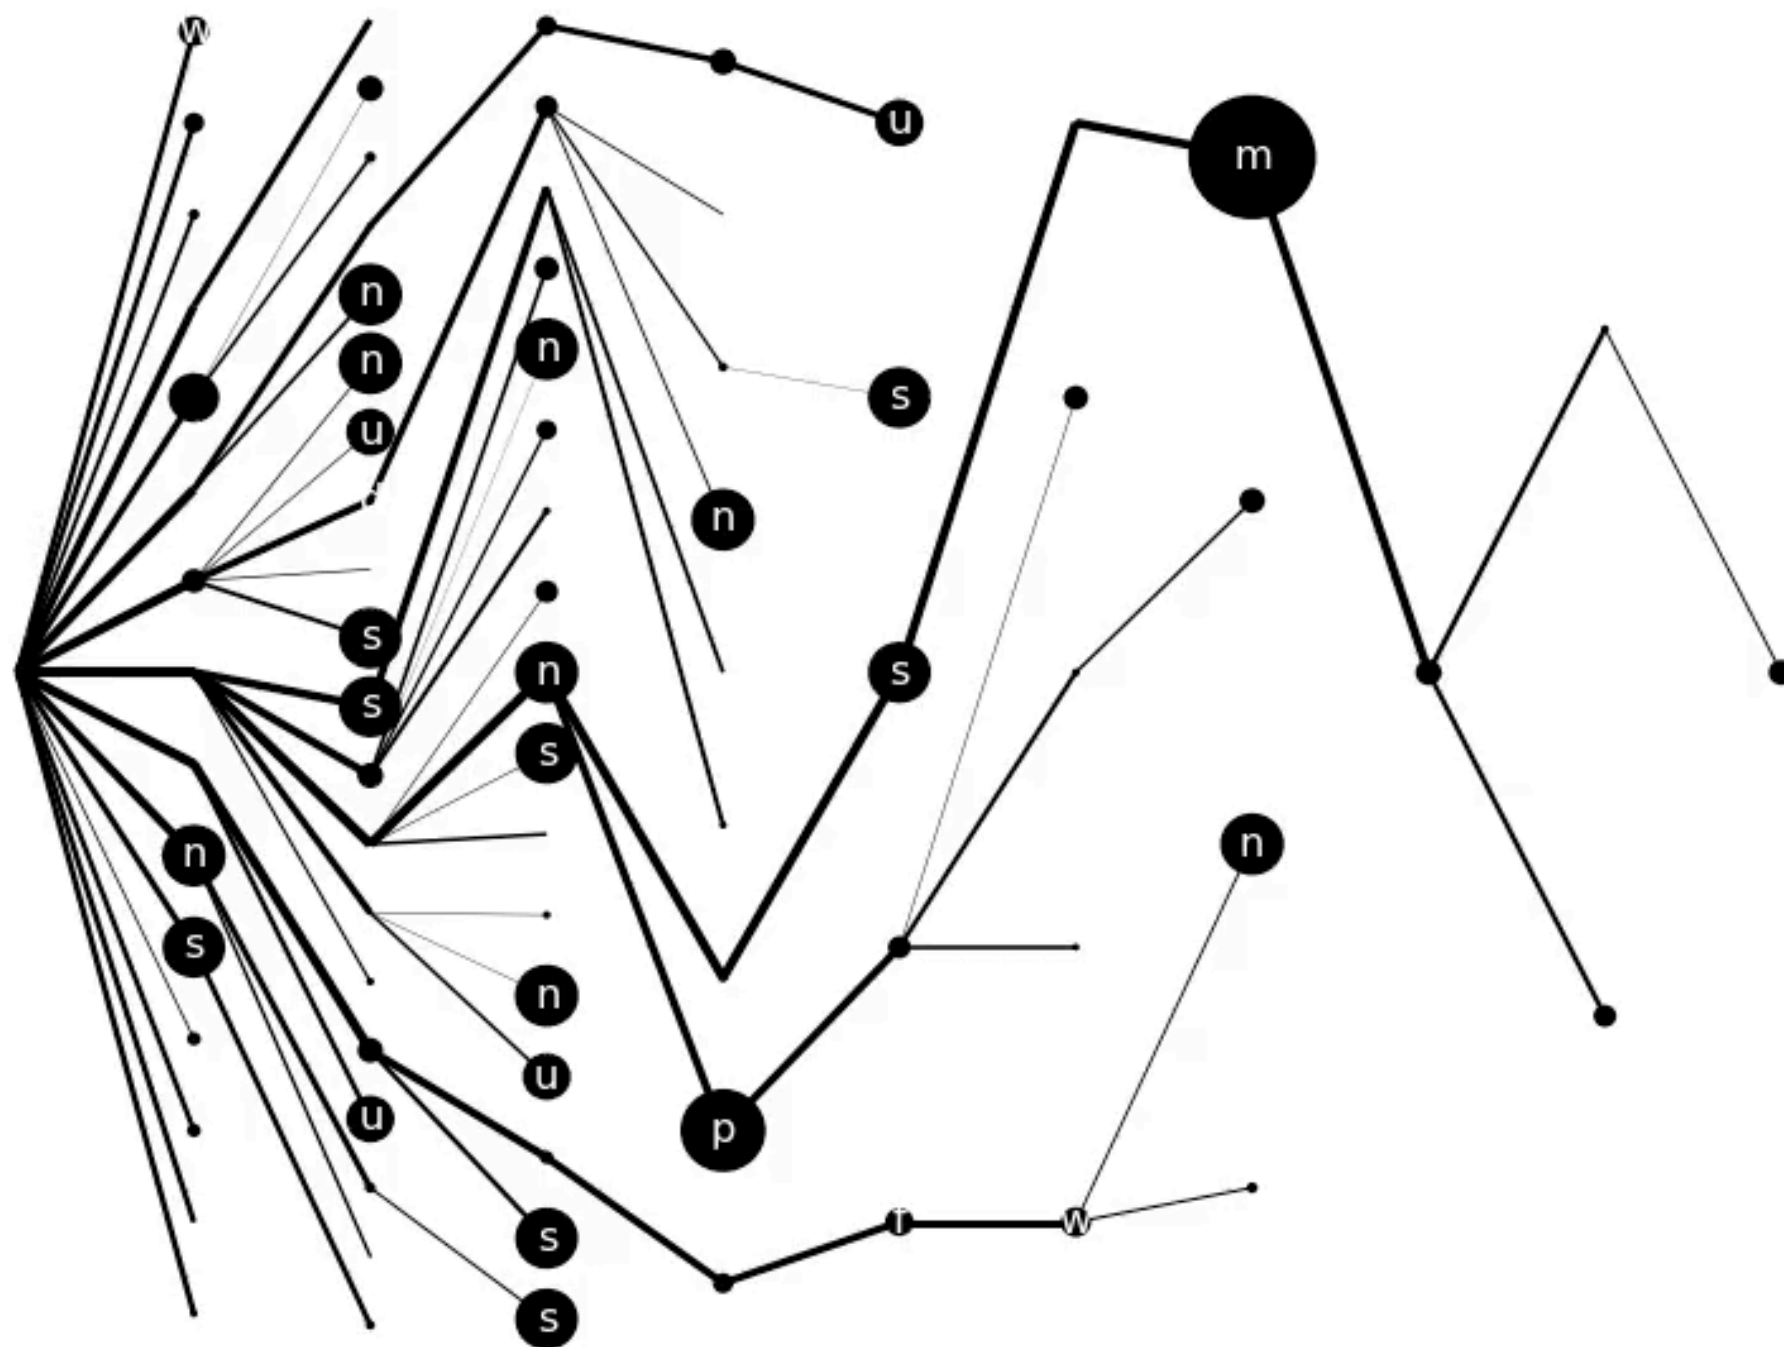
- Start with 20 top packages (SOCK paper)
- Eliminate 2 that don't work with OpenLambda currently
- Add 14 dependencies


1000 Synthetic Lambdas

- For each, pick 2 of the 32 packages, uniformly at random (the 2 picks are independent)
- Import those 2 (this usually entails importing deps too)
- By importing a package, we mean importing its top-level modules
- Do nothing else (no-op)

Client:

- Invoke each Lambda once
- 5 concurrent invocations
- Measure throughput





Observation: up to ~40 Zygote nodes are useful, and gives us ~16 invocations/sec (this is on a 2-core, 4-GB VM on my Macbook)

More Methodology Details

(previous and subsequent plots)



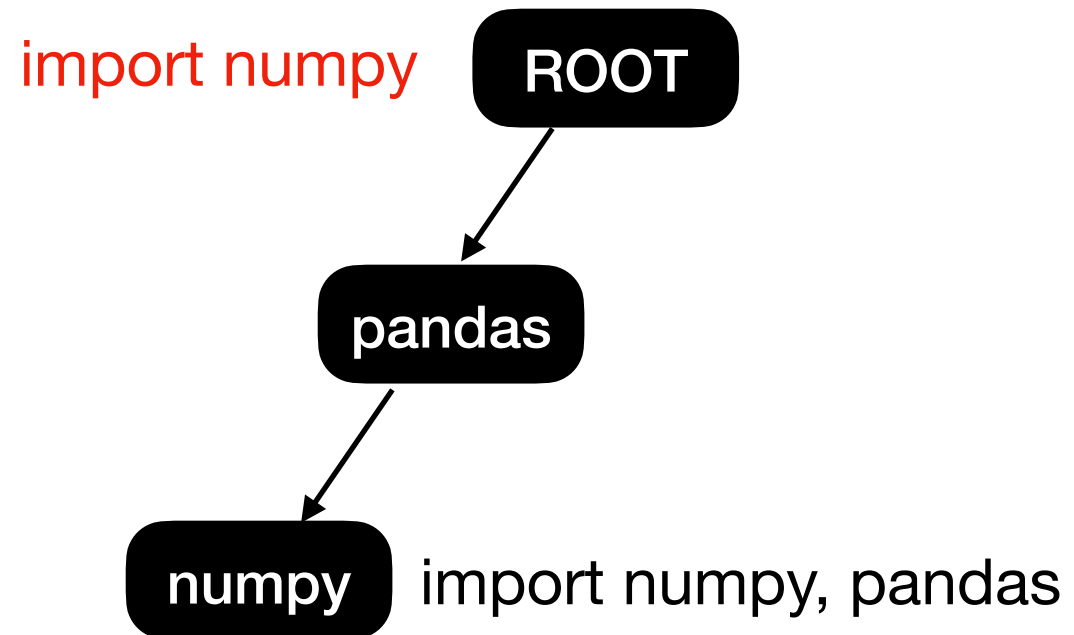
Overfitting


- Generate pairs of train/test traces
- Fit Zygotree to **train trace**
- Measure throughput of that tree on **test trace**

Variance

- Generate 5 train/test pairs
- Gray regions represent standard deviation
- Each test trace is executed for various tree sizes, so the same sample of 5 is reused along the x-axis


Rule I: Put Prereqs Closer to Root






Observation: enforcing pre-reqs didn't help, perhaps because the workload tends to make prereqs be used more often (so they'll usually be near the root anyway)


Rule 2: Put Heavy Modules Near Root



Assume there are no package deps between A and B

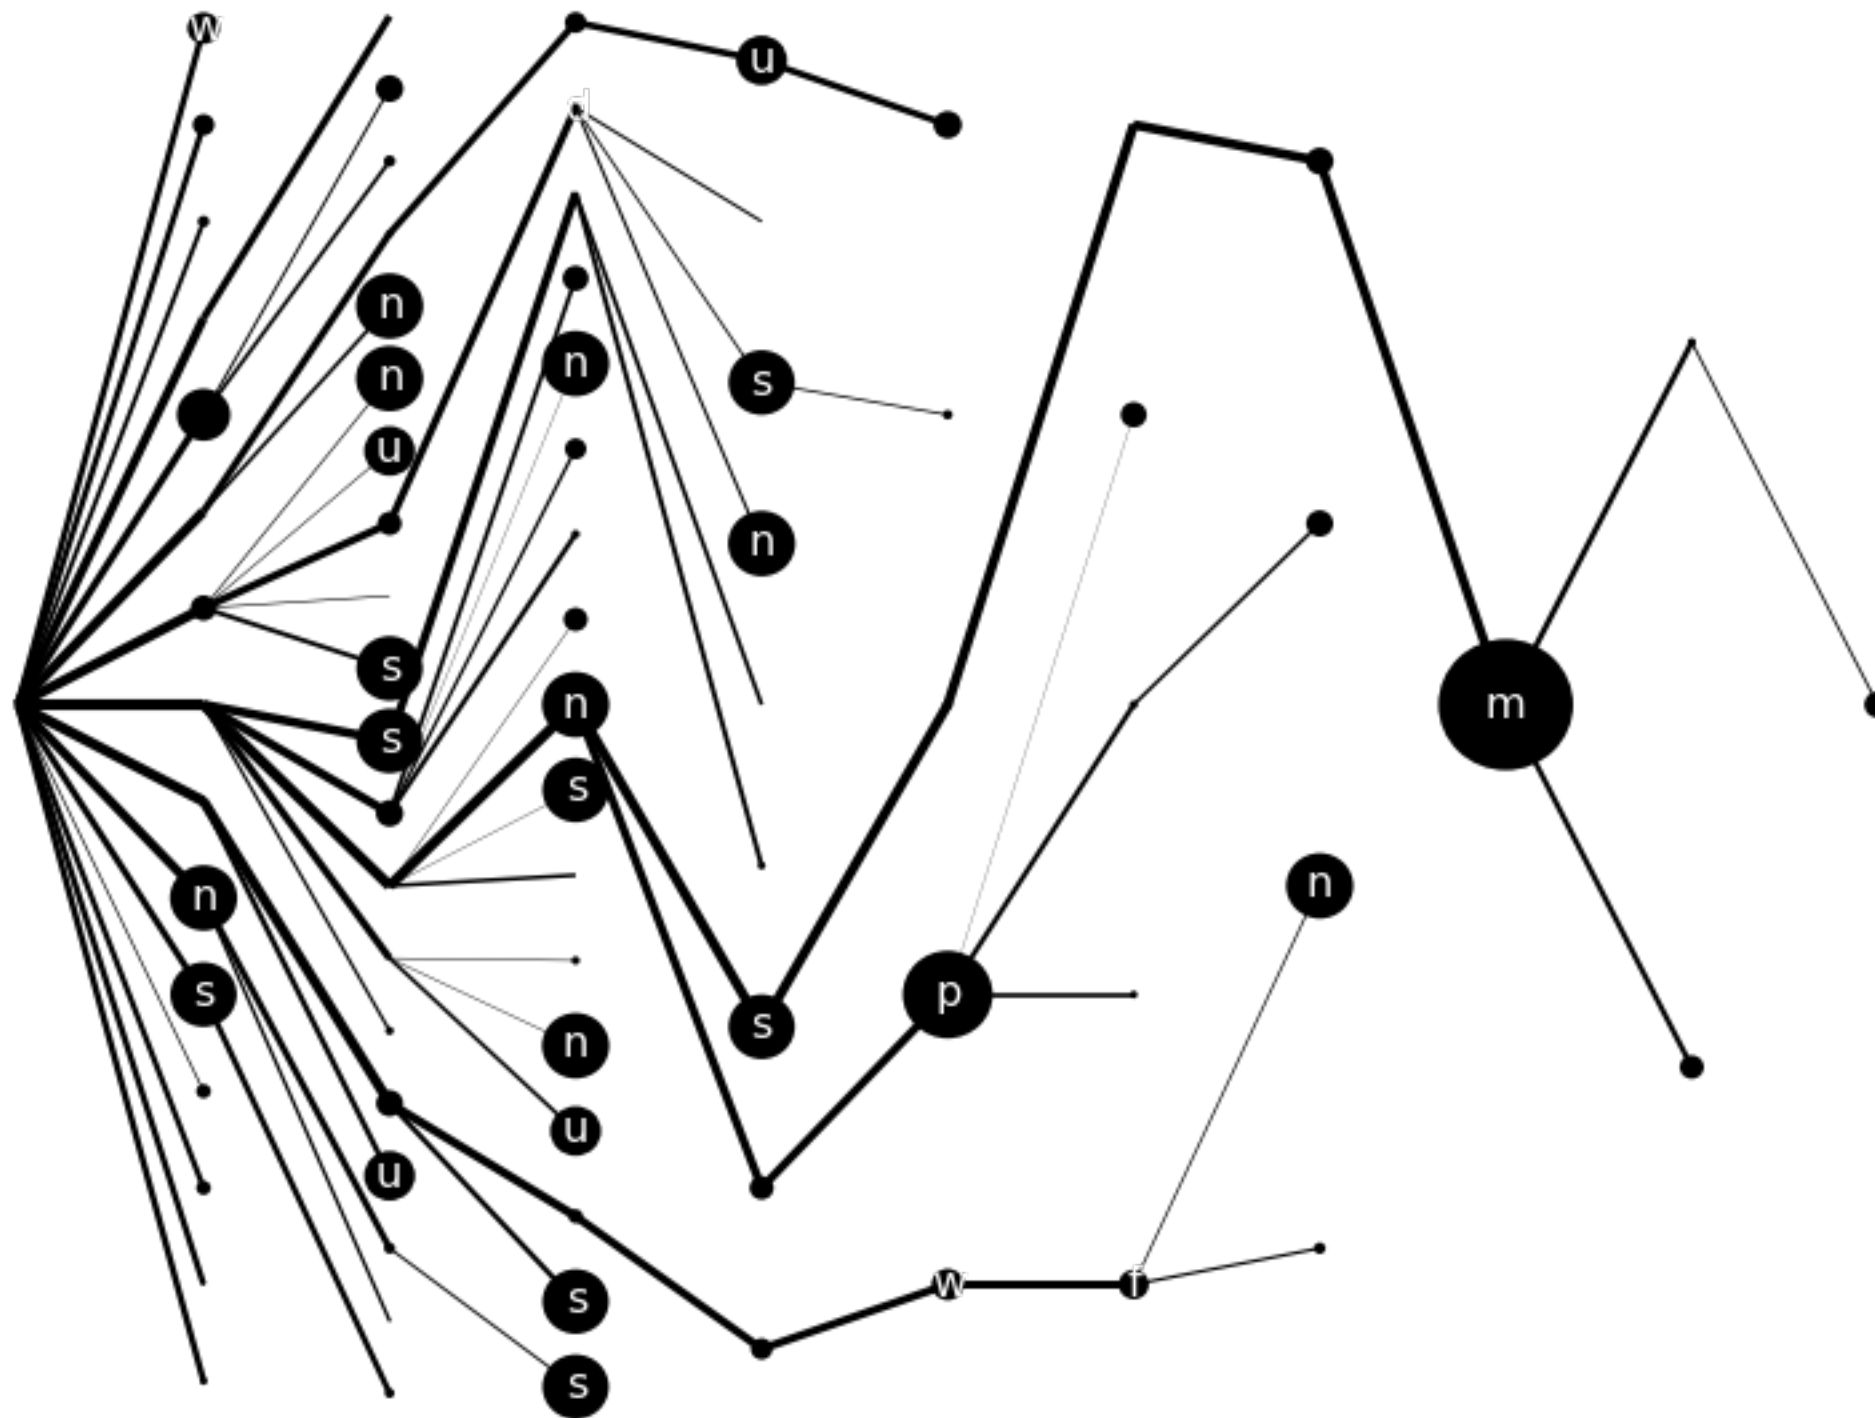


We use **import time** as the weight/cost of a package
(this is somewhat correlated with the alternative metric, mem footprint)




Observation: this rule really hurts perf! Why?


Without Rule 2



With Rule 2




With Rule 2




Observation: the algorithm is greedy. We're not getting to very heavy packages (e.g., pandas + matplotlib) that depend on very light packages.

Rule 3: Split to Minimize Entropy




	pkg A	pkg B	pkg C	pkg D
fn 1	1	1	0	1
fn 2	1	0	1	0
fn 3	0	1	0	1
fn 4	0	0	1	0




Observation: this addresses one kind of shortsighted greediness, but doesn't solve the priority inversion problem.

Rule 4: Distribute weights to dependencies (like "priority inheritance")




A, C, and D have no deps. B is large and depends on A, which is tiny.

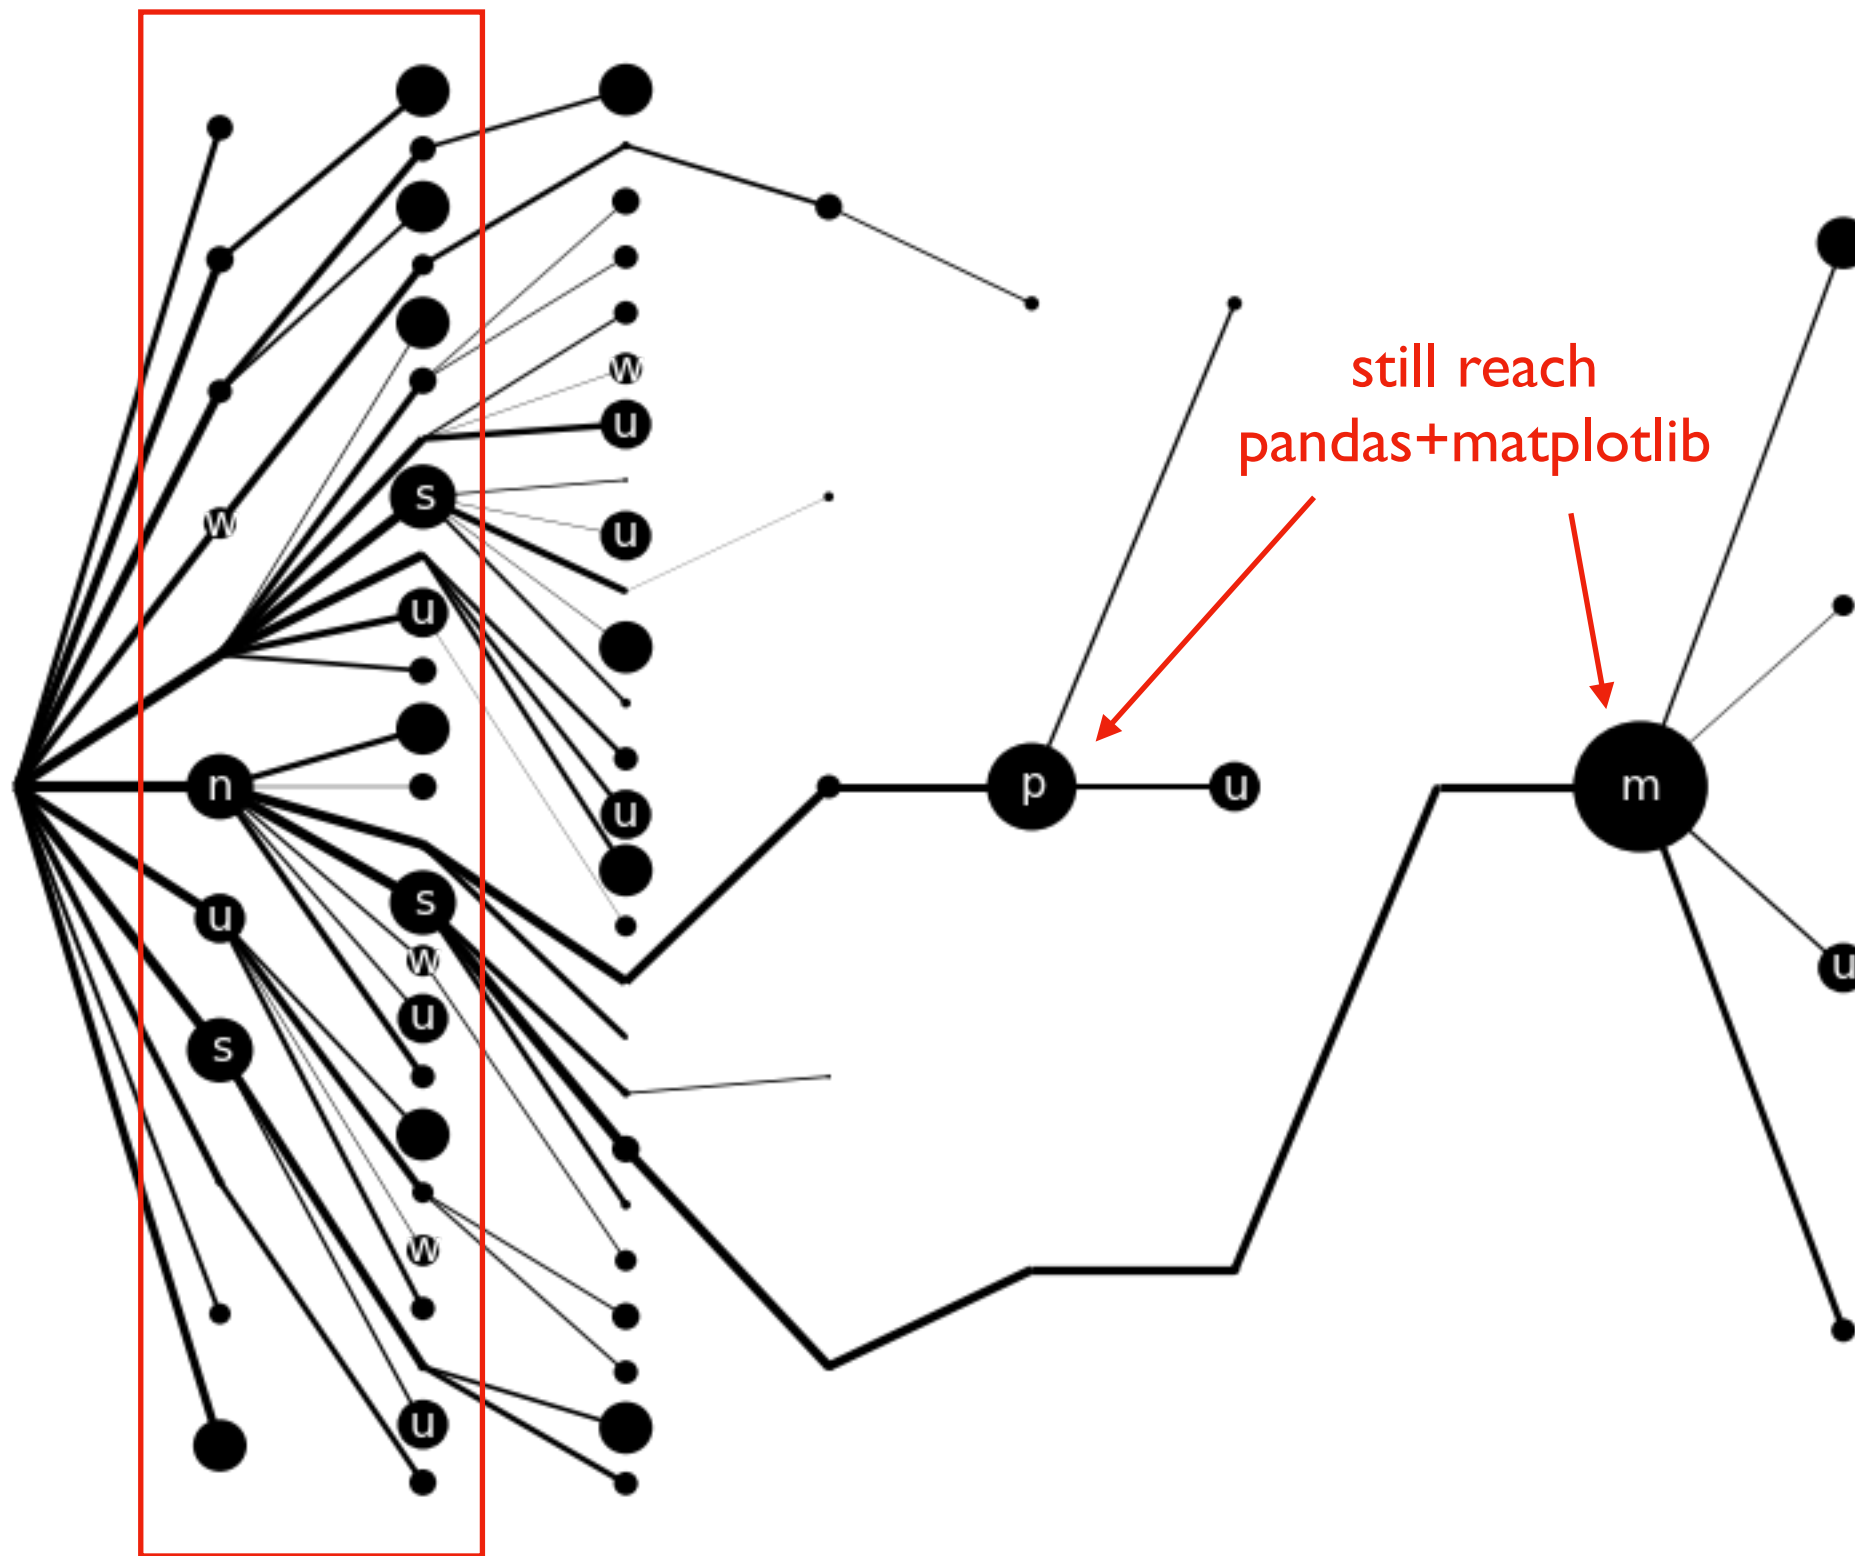
Without Rule 4



With Rule 4




With Rule 4



lots of heavy packages

still reach
pandas+matplotlib



Observation: adding priority inheritance gives best performance!

Concluding Thoughts...

1. Start with a smart static policy

- Make dynamic as a tweak

2. What ML algorithms can we repurpose as policies?

- A small tweak to the decision tree algorithm converts it from a classifier to a cache policy